

x86 and PC architecture

PC architecture

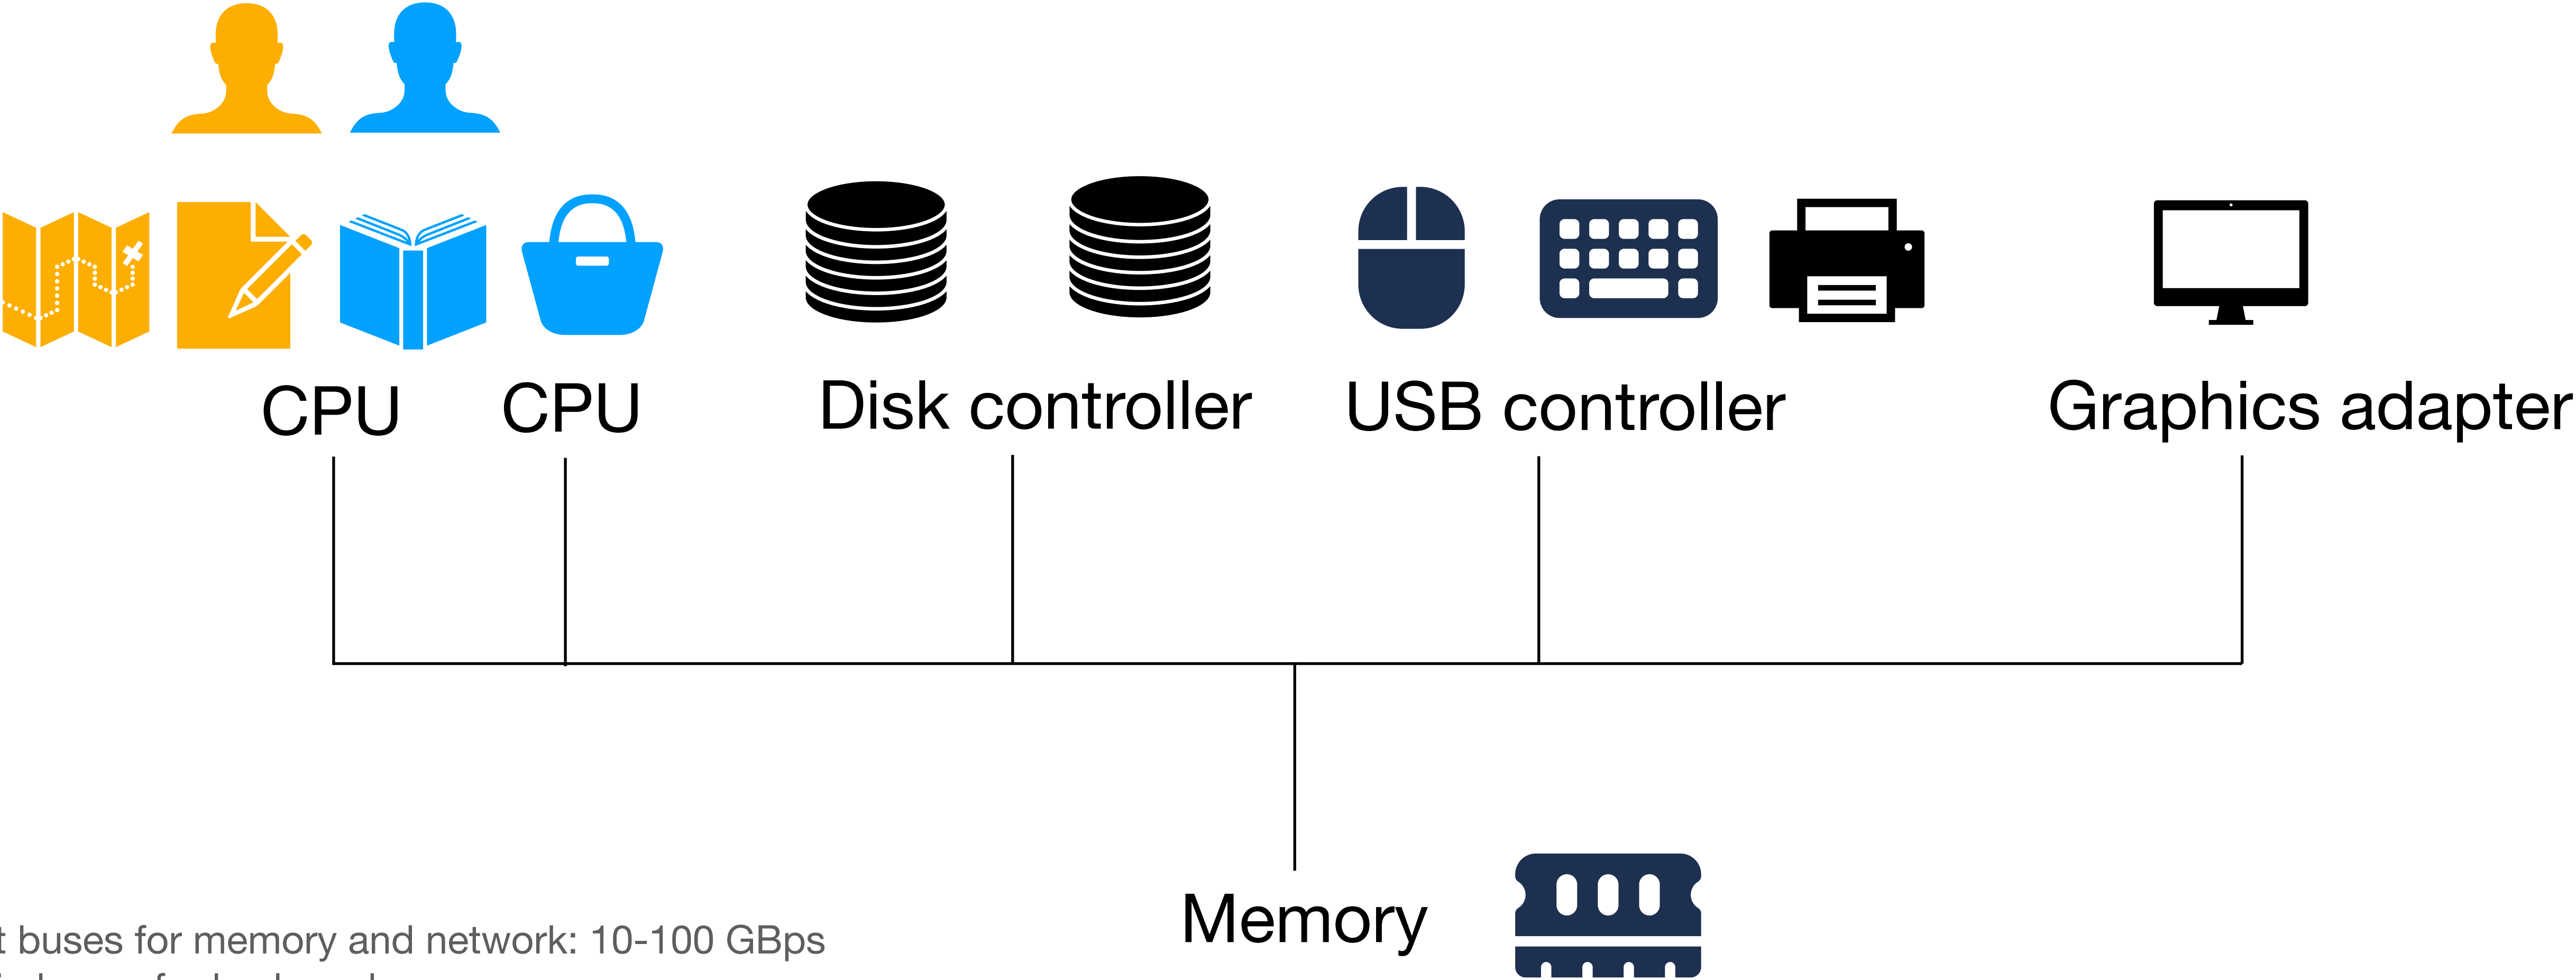
x86 instruction set

gcc calling convention

Agenda

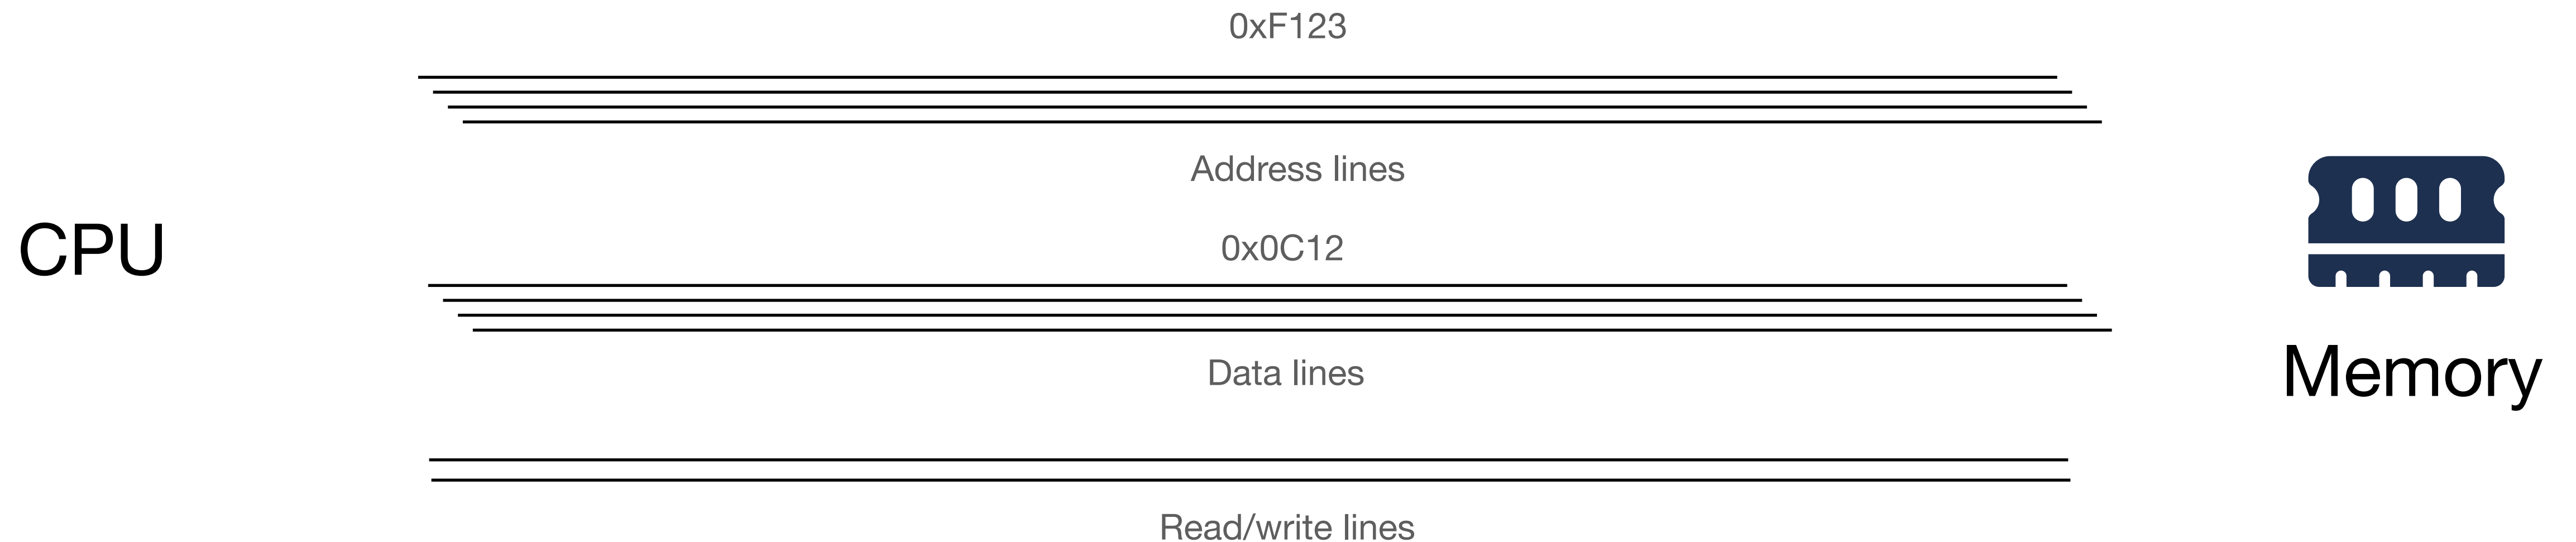
- Build a mental model of how PC components (e.g., CPU and memory) interact with one another
- x86 instruction set: Defined by Intel in early 1980s. Has become a standard. Understand x86 instruction set so that we can read and write x86 assembly
 - Assembly programs are sometimes required by OS to get fine-grained control of the hardware
- Understand gcc calling convention so that we can call C programs from assembly and vice-versa

Computer organization



Fat buses for memory and network: 10-100 GBps
Thin buses for keyboard, mouse

CPU-memory interaction



- Each read/write takes ~100 cycles
- Faster memory: on-chip registers ~1 cycle.

Registers

- **General purpose registers.**
 - **%eax, %ebx, %ecx, %edx**
 - **%edi: destination index, %esi: source index**
- Flags register. %eflags
- Instruction pointer. %eip
- Stack registers. %ebp: base pointer, %esp: stack pointer
- Special registers.
 - Control registers %cr0, %cr2, %cr3, %cr4;
 - Segment registers %cs, %ds, %es, %fs, %gs, %ss
 - Global and local descriptor table registers %gdtr, %ldtr
- Other registers not used in xv6: 8 80-bit floating point registers, debug registers

mov instructions

Intel SDM Vol 1 7.3.1.1

Assembly	“C” equivalent
movl %eax, %edx	edx = eax
movl \$123, %edx	edx=0x123
movl 0x123, %edx	%edx = *(int32_t*)0x123
movl (%ebx), %edx	edx=*(int32_t*) ebx
movl 4(%ebx), %edx	edx=*(int32_t*)(ebx+4)

Assembly	“C” equivalent
movsb	*edi = *esi; edi++; esi++;

Other instruction variants

General-Purpose Registers						
31	16	15	8	7	0	
			AH		AL	16-bit
			BH		BL	AX
			CH		CL	BX
			DH		DL	CX
			BP			DX
			SI			32-bit
			DI			EAX
			SP			EBX
						ECX
						EDX
						EBP
						ESI
						EDI
						ESP

Figure 3-5. Alternate General-Purpose Register Names

- movw: moves 2 bytes (%ax)
- movb: moves 1 byte (%al, %ah)

Many other instructions: ADD, SUB, MUL, DIV, ...

Registers

- General purpose registers.
 - `%eax`, `%ebx`, `%ecx`, `%edx`
 - `%edi`: destination index, `%esi`: source index
- **Flags register. `%eflags`**
- Instruction pointer. `%eip`
- Stack registers. `%ebp`: base pointer, `%esp`: stack pointer
- Special registers.
 - Control registers `%cr0`, `%cr2`, `%cr3`, `%cr4`;
 - Segment registers `%cs`, `%ds`, `%es`, `%fs`, `%gs`, `%ss`
 - Global and local descriptor table registers `%gdtr`, `%ldtr`
- Other registers not used in xv6: 8 80-bit floating point registers, debug registers

EFLAGS

- Carry flag: Most significant bit overflowed.

```
movl $FFFFFFFF %eax
addl %eax, %eax
```

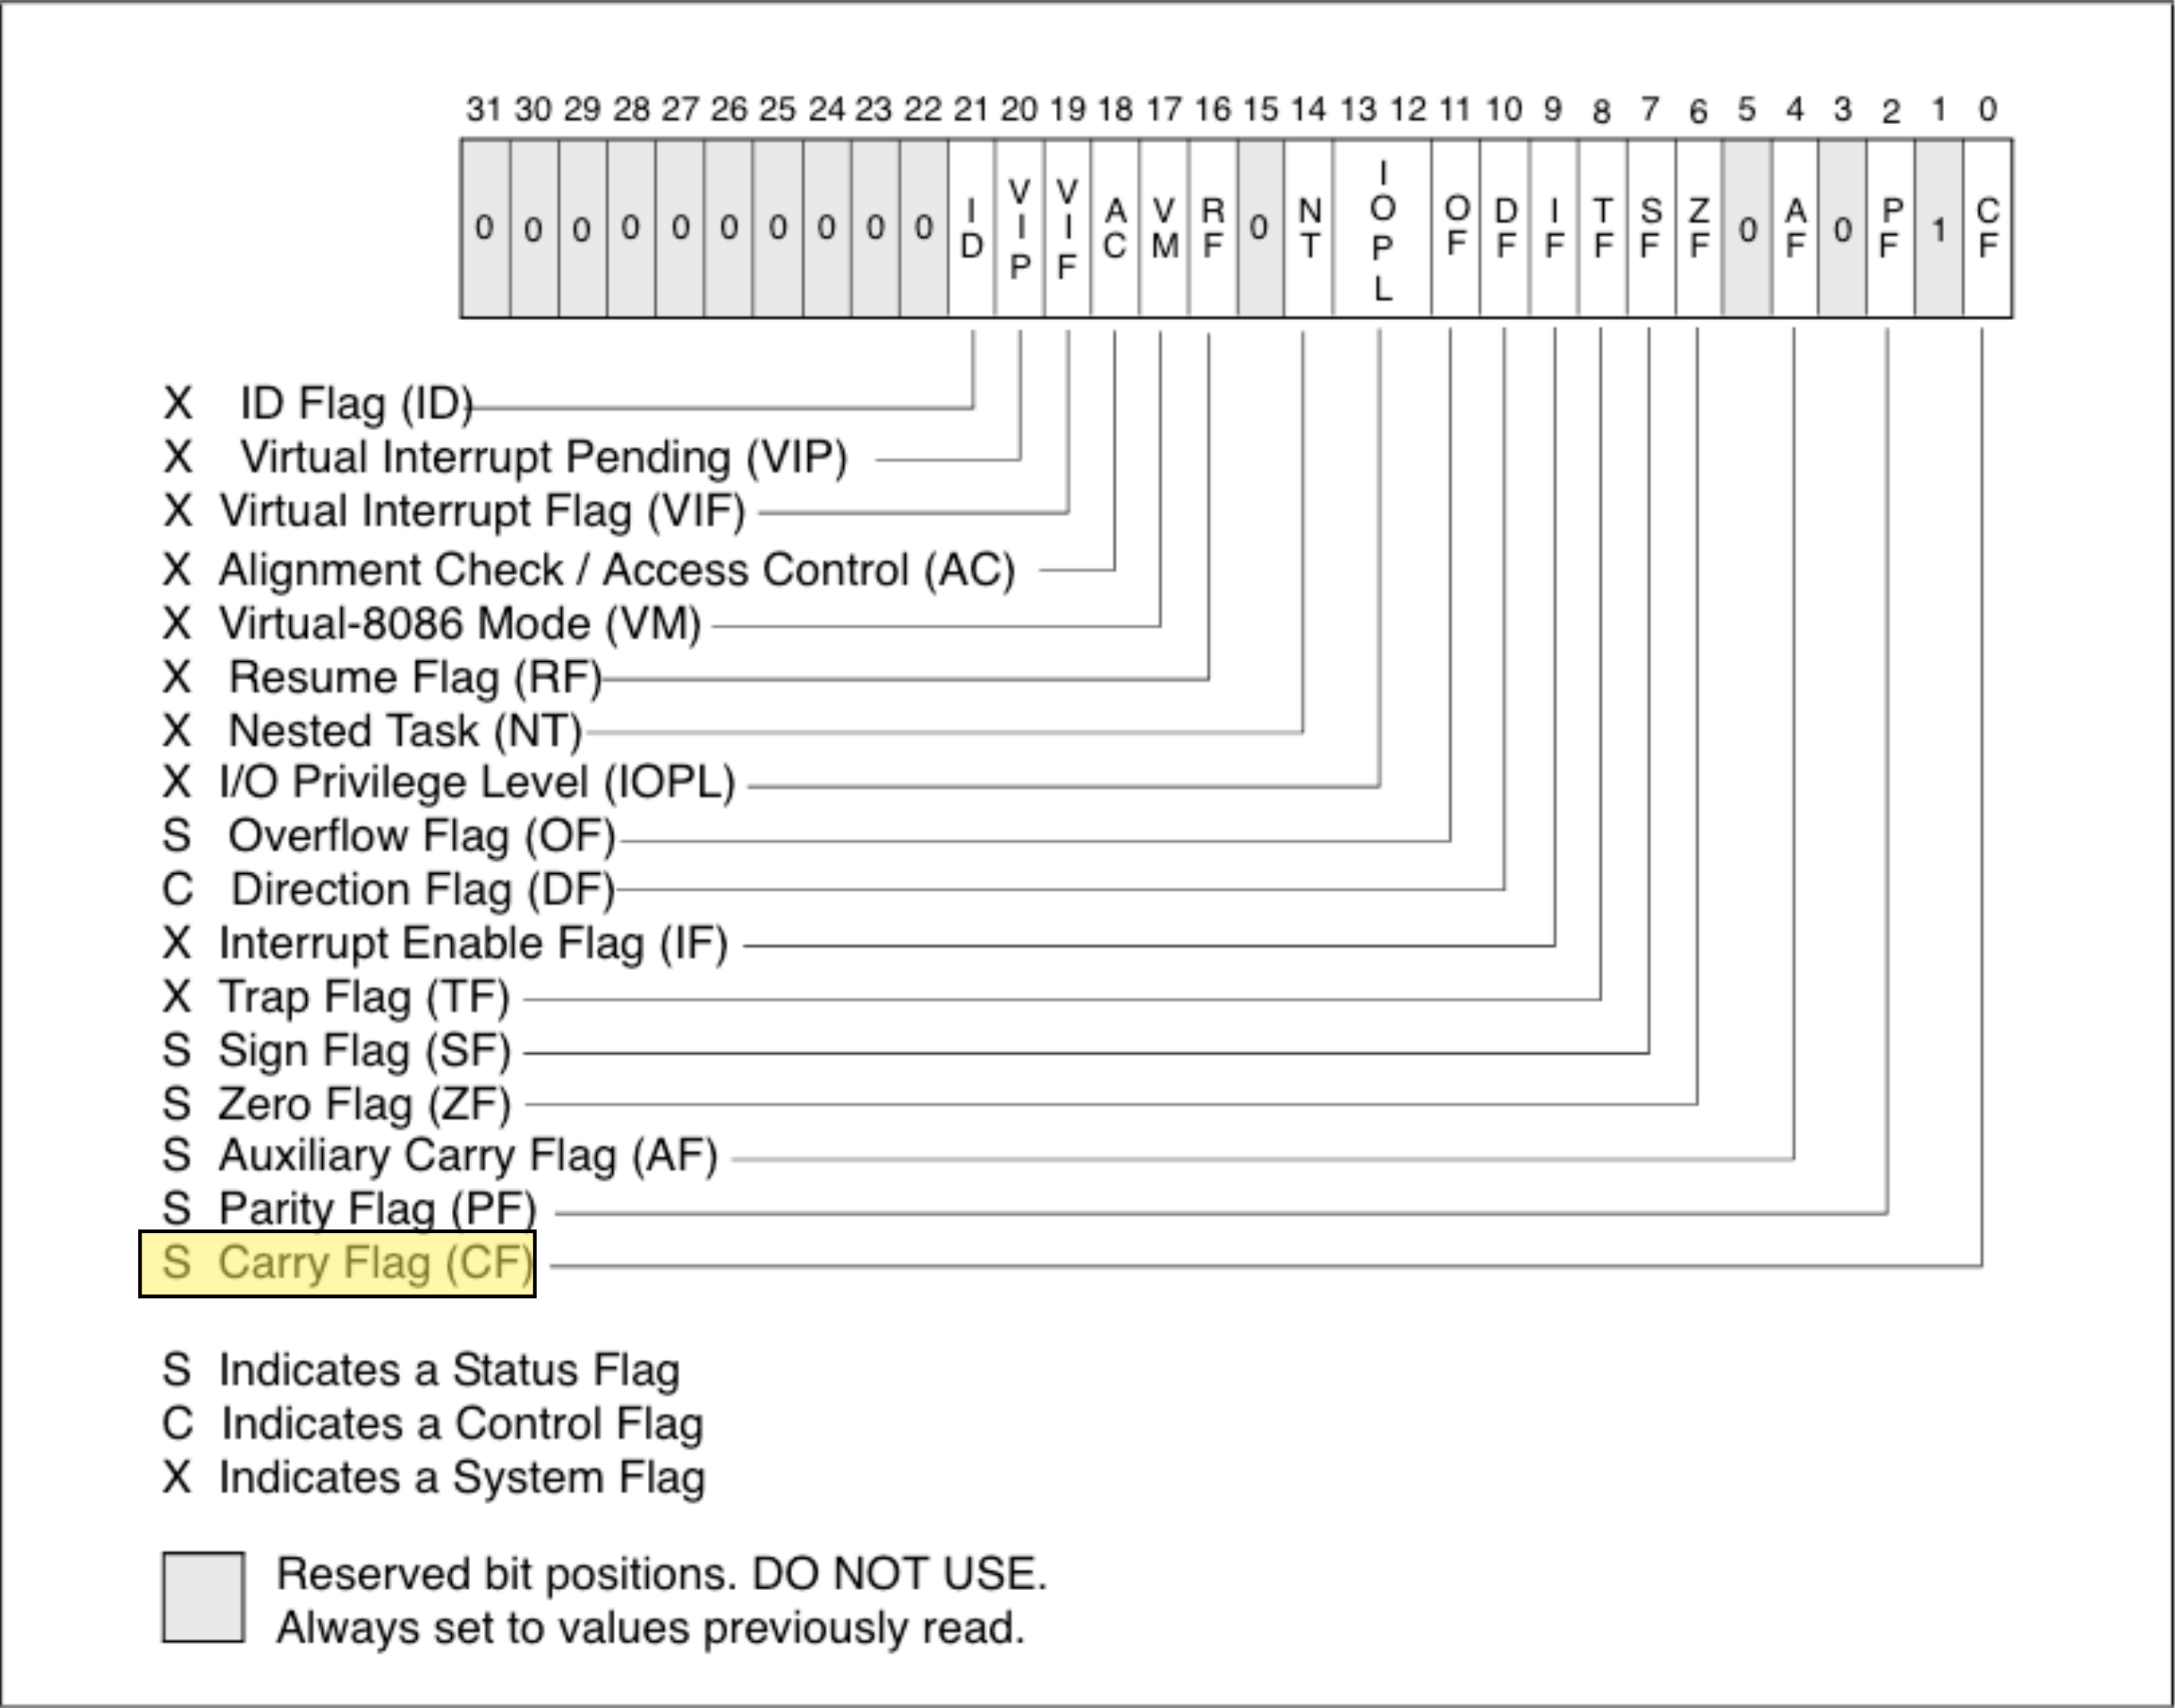


Figure 3-8. EFLAGS Register

EFLAGS (2)

- Zero flag: Set if result is zero.

```
xorl %eax, %eax
```

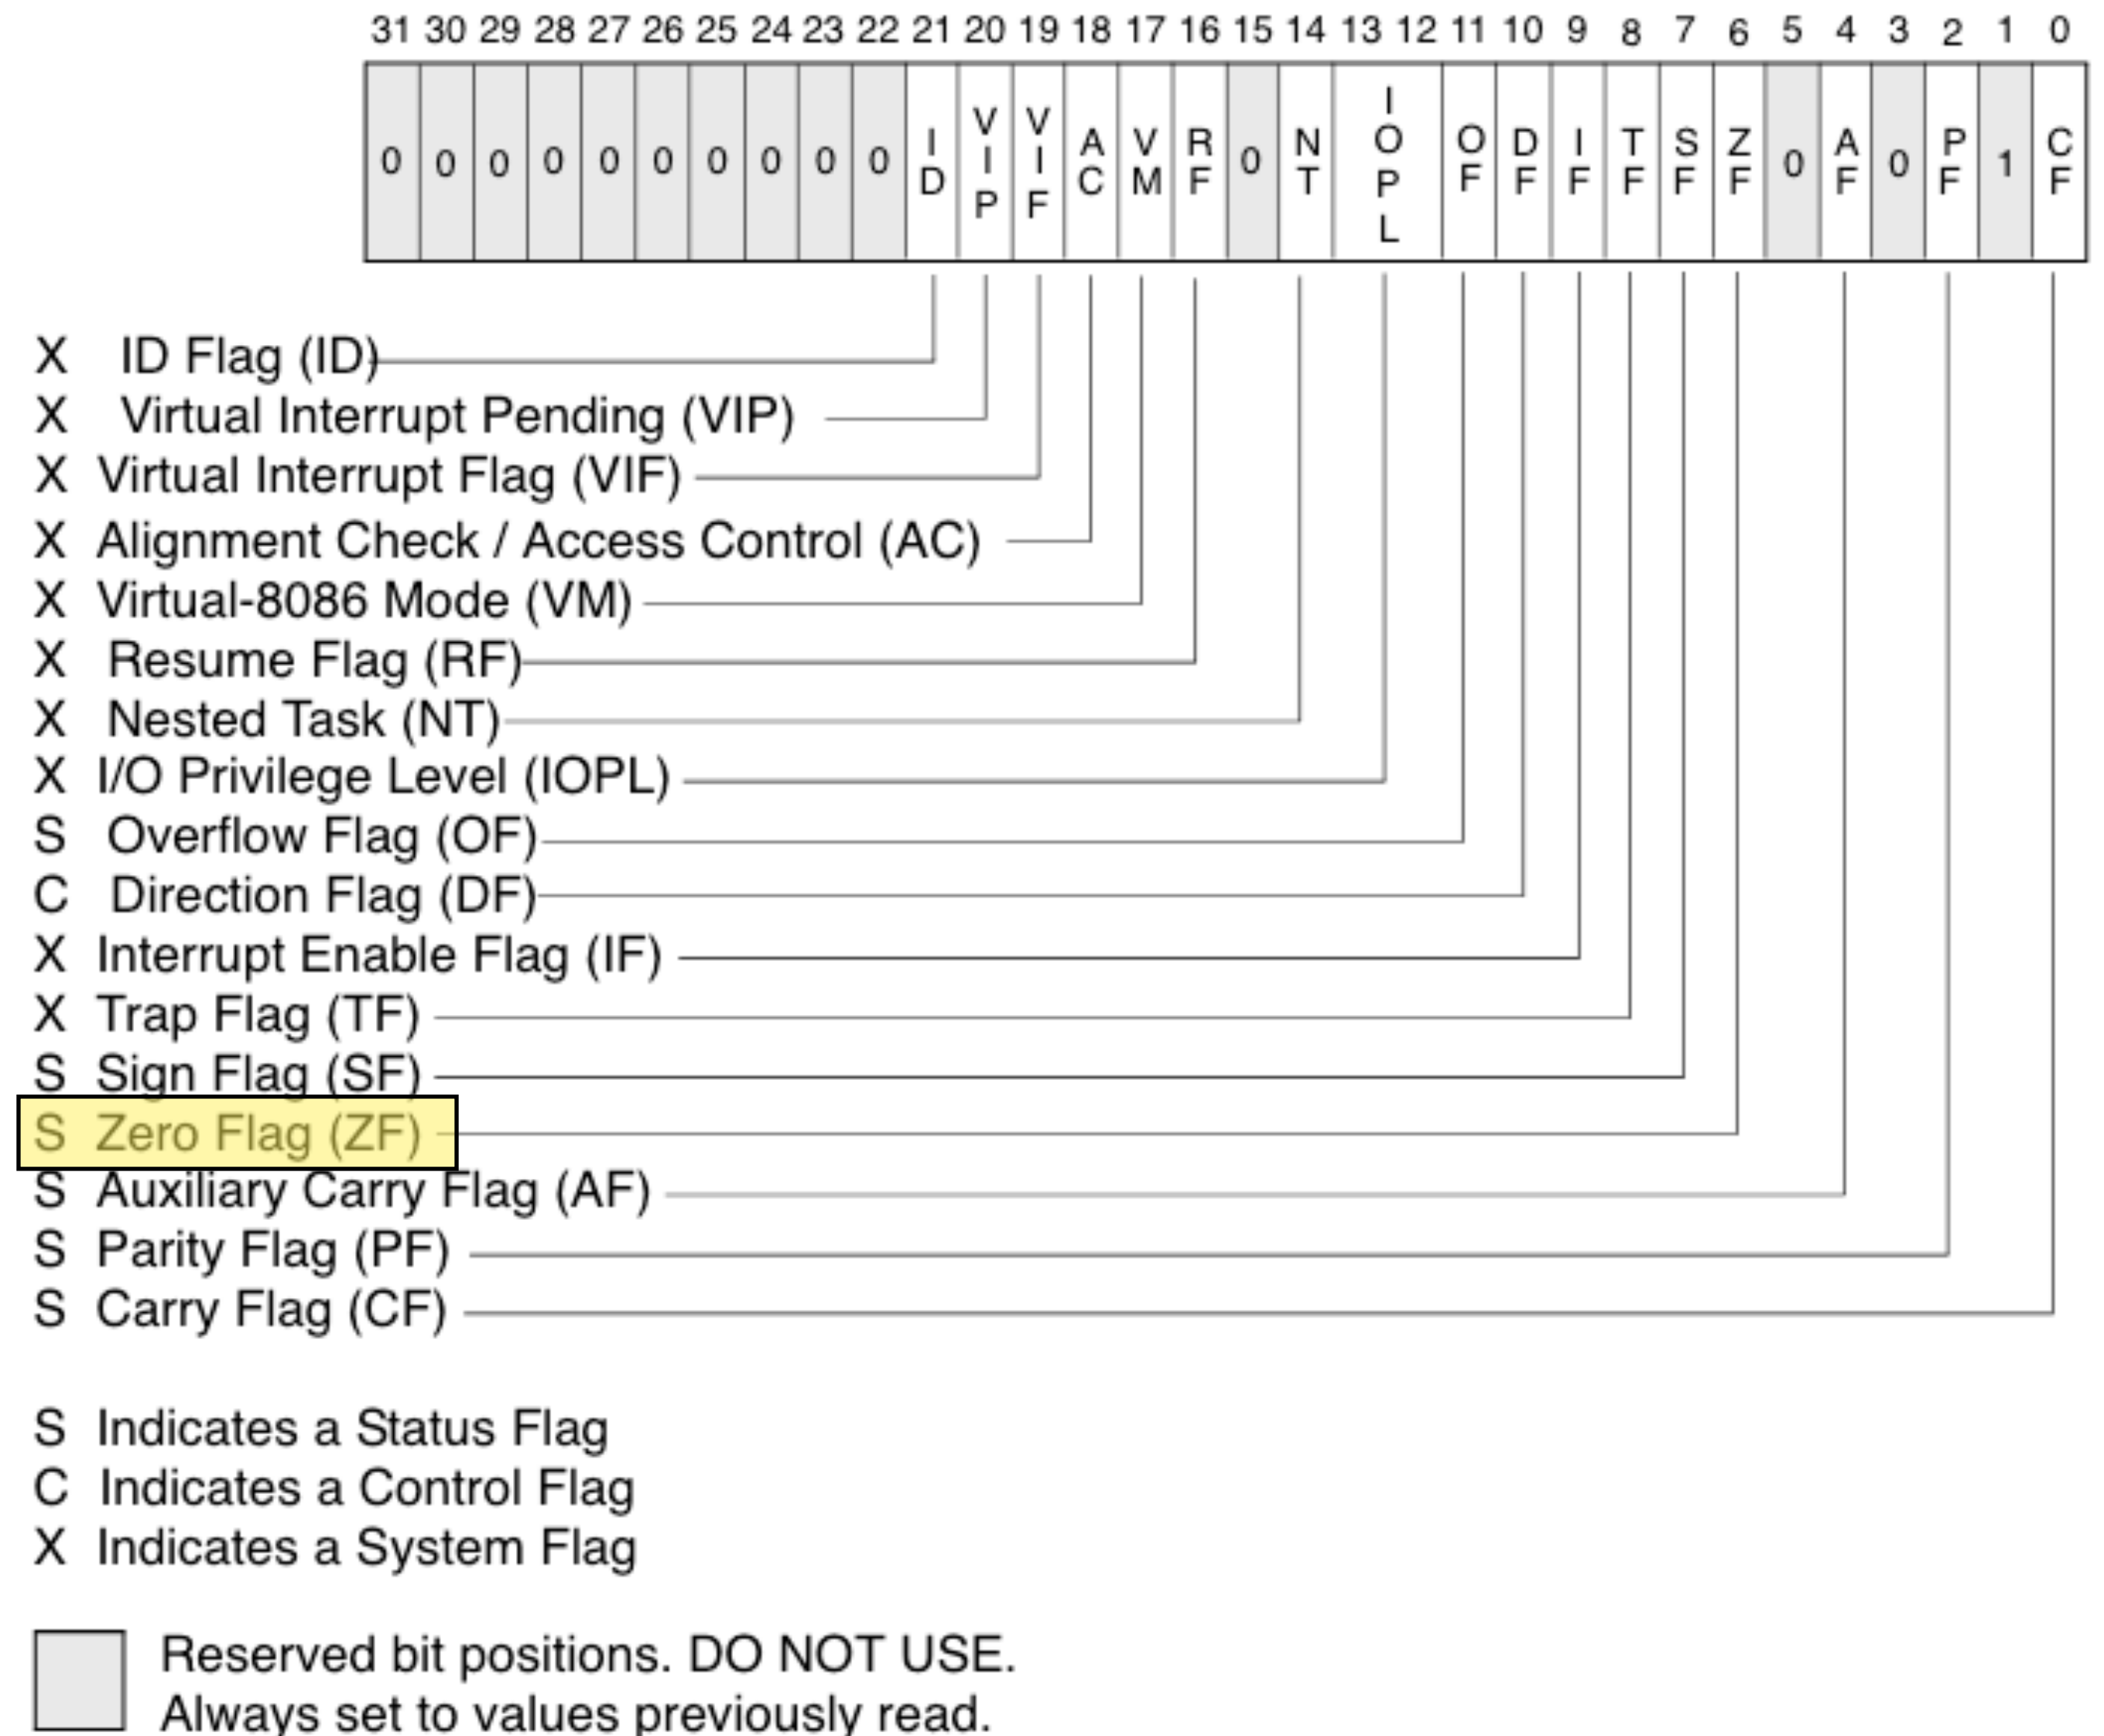


Figure 3-8. EFLAGS Register

EFLAGS (3)

- Sign flag: Equal to the most significant bit of the result (which is the sign bit of a signed integer)

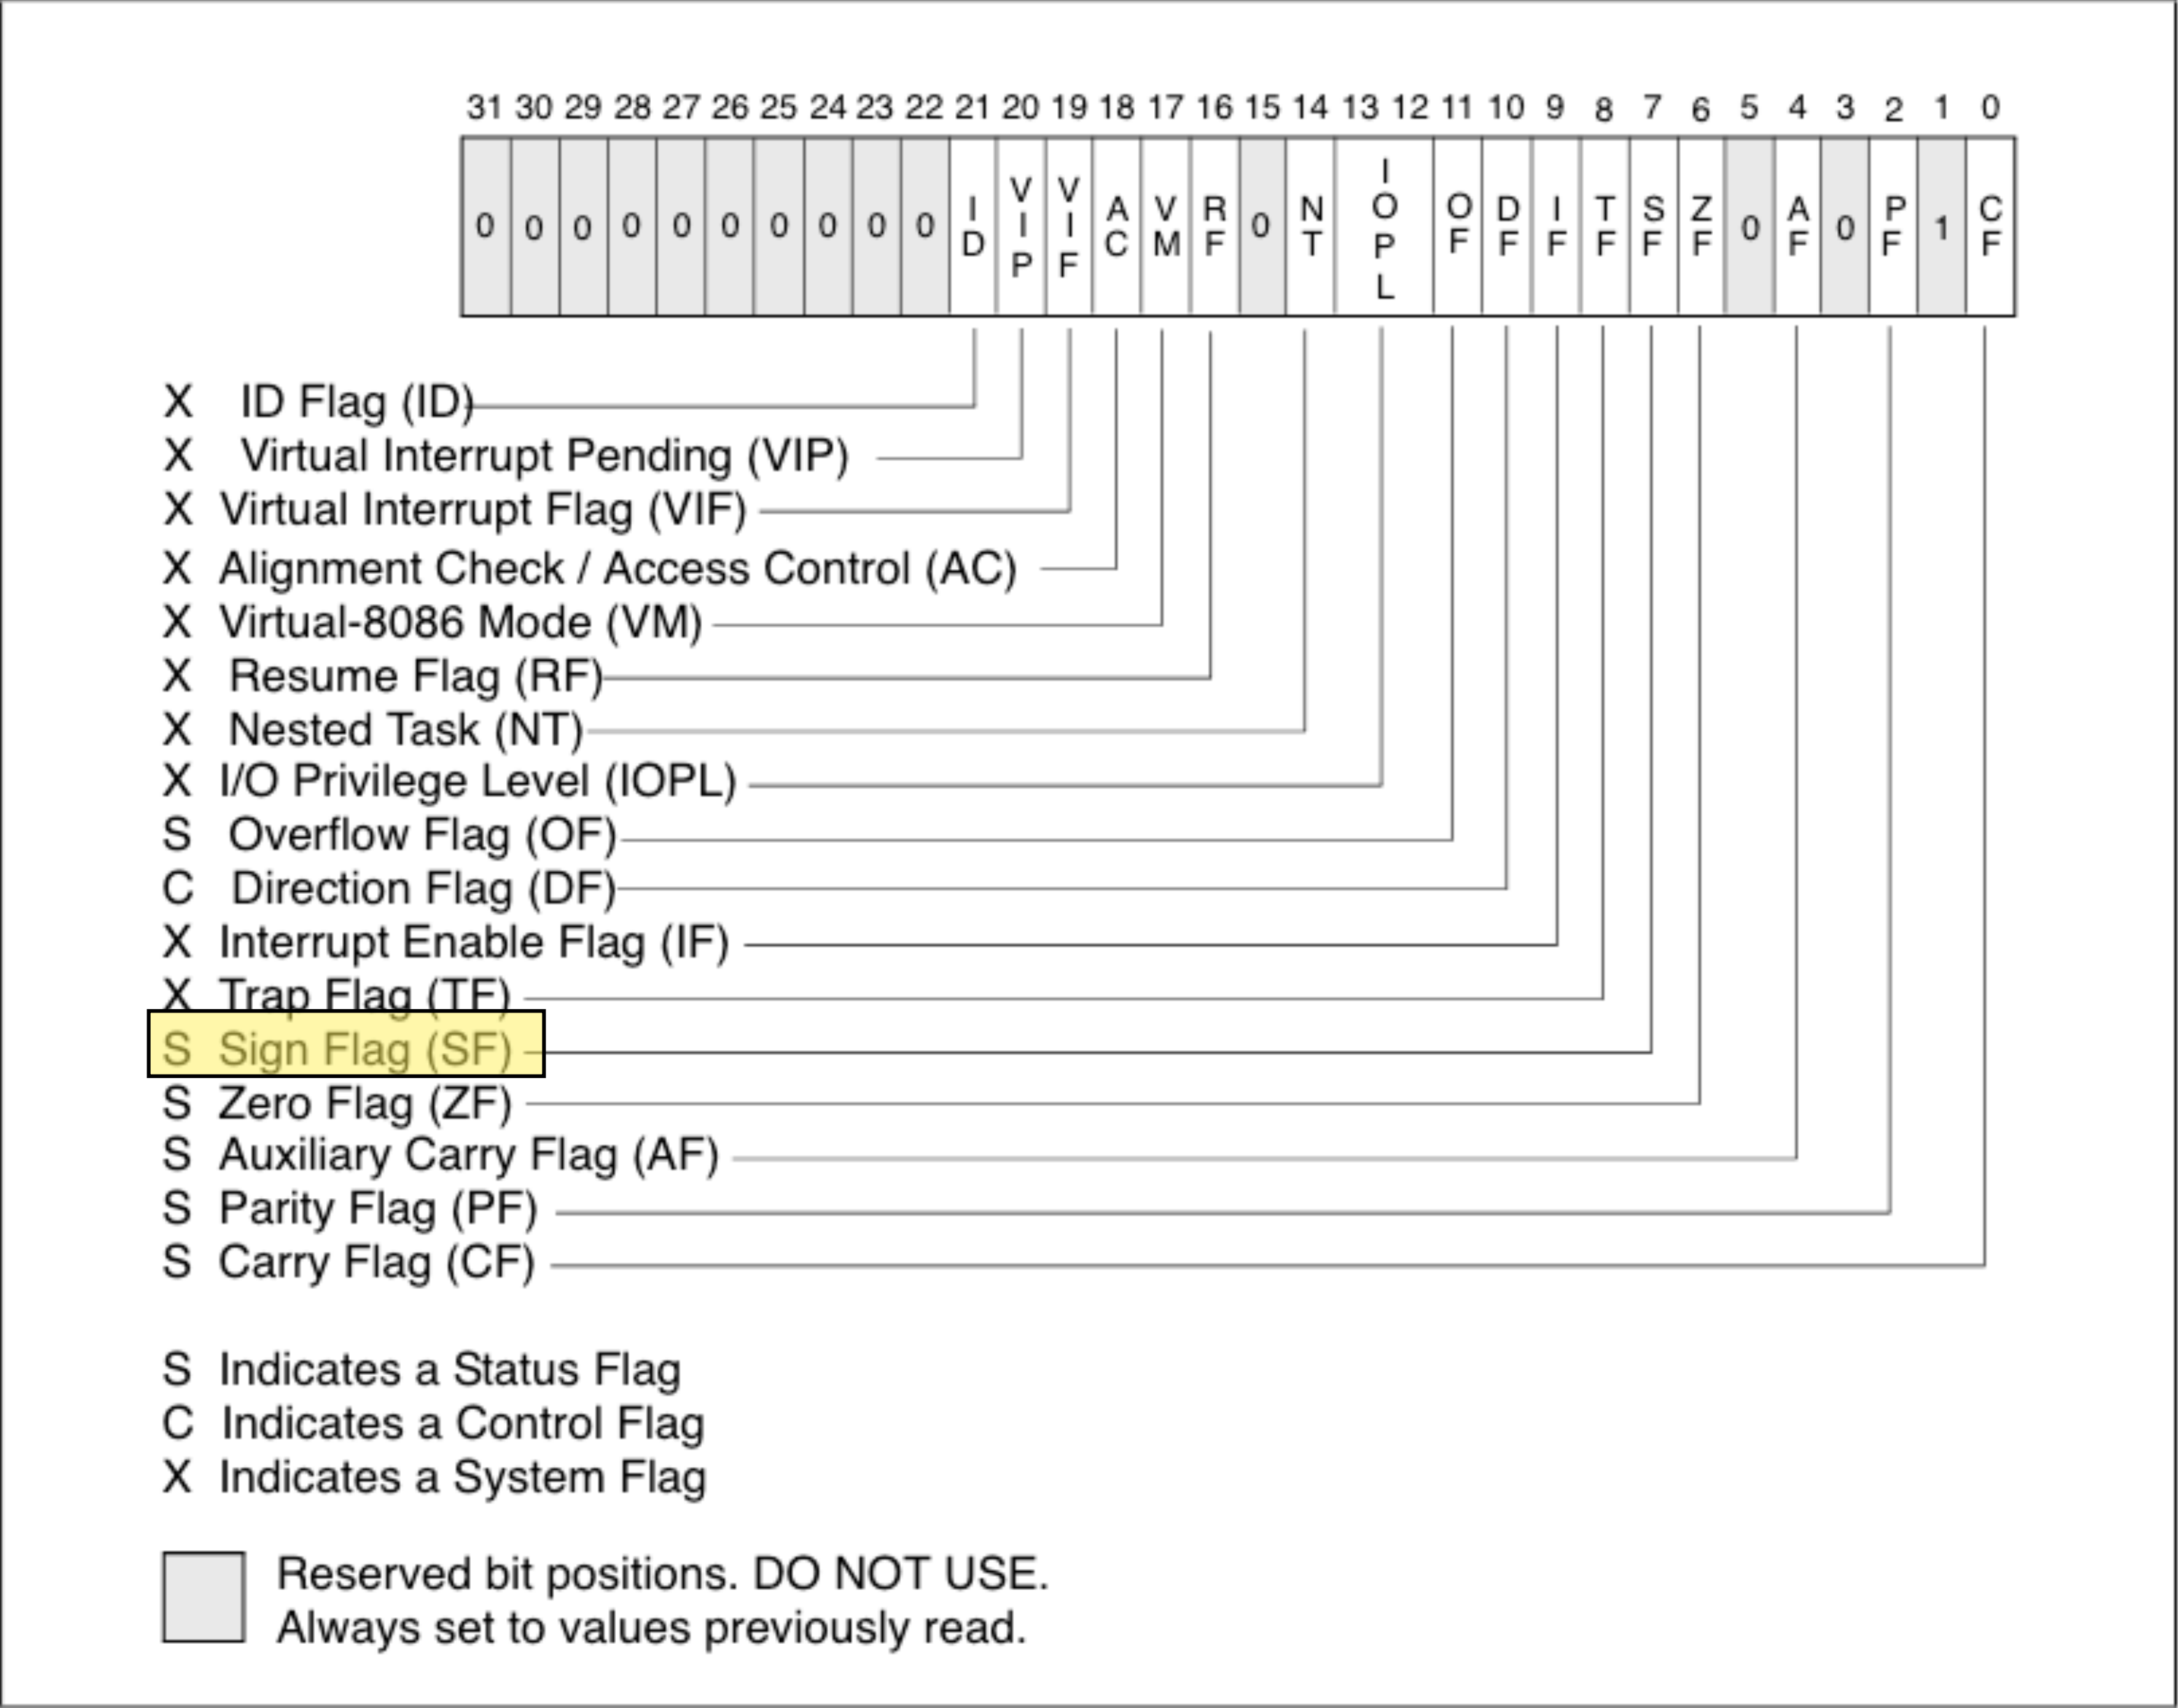



Figure 3-8. EFLAGS Register

Registers in action

02.flags.c

```
int foo(int x, int y) {  
    int z = x + y;  
    if(z % 2 == 0)  
        return x;  
    return y;  
}
```

gcc -m32 -S -O1 02.flags.c



02.flags.s

```
foo:  
    movl 4(%esp), %eax    # eax = x  
    movl %eax, %edx      # edx = eax (z = x)  
    addl 8(%esp), %edx    # edx += y  
    andl $1, %edx        # edx = (edx & 1). ZF if edx is even.  
    cmovne 8(%esp), %eax  # eax = y if !ZF  
    ret
```

Registers

- General purpose registers.
 - `%eax`, `%ebx`, `%ecx`, `%edx`
 - `%edi`: destination index, `%esi`: source index
- Flags register. `%eflags`
- **Instruction pointer. `%eip`**
- Stack registers. `%ebp`: base pointer, `%esp`: stack pointer
- Special registers.
 - Control registers `%cr0`, `%cr2`, `%cr3`, `%cr4`;
 - Segment registers `%cs`, `%ds`, `%es`, `%fs`, `%gs`, `%ss`
 - Global and local descriptor table registers `%gdtr`, `%ldtr`
- Other registers not used in xv6: 8 80-bit floating point registers, debug registers

Instruction pointer

- Next instruction is pointed to by instruction pointer %eip

```
for(;;){  
    run next instruction  
}
```

- %eip is simply incremented in most cases
- Except special instructions
 - JMP 0x1234: changes %eip to 0x1234 e.g., while loop
 - JP, JN, J[N]Z: jump if last result was positive, negative, zero, non-zero etc. This uses bits from EFLAGS register. e.g, if(x > 0) { .. }
 - CALL 0x1234: Similar to JMP, additionally saves the current instruction pointer on stack e.g., function call
 - RET: returns back to callee. Changes %eip to address in stack

Registers in action (2)

02.eip.c

```
int exponent(int x, int y) {  
    int z = x;  
    while(y > 0) {  
        z = z * x;  
        y --;  
    }  
    return z;  
}
```

gcc -m32 -S -O1 02.eip.c

exponent:

```
    movl 4(%esp), %ecx  
    movl 8(%esp), %eax  
    movl %ecx, %edx  
    testl %eax, %eax
```

```
    jle .L1
```

.L3:

```
    imull %ecx, %edx  
    subl $1, %eax  
    jne .L3
```

.L1:

```
    movl %edx, %eax  
    ret
```

ecx = x

eax = y

edx = ecx (z = x)

bitwise and eax with eax.

SF if eax<0. ZF if eax=0.

Jump if SF or ZF (y <= 0)

z = z*x

eax-- (y--). ZF if eax=0 (y=0)

Jump back to loop if !ZF

eax = edx (return z)

Registers

- General purpose registers.
 - `%eax`, `%ebx`, `%ecx`, `%edx`
 - `%edi`: destination index, `%esi`: source index
- Flags register. `%eflags`
- Instruction pointer. `%eip`
- **Stack registers. `%ebp`: base pointer, `%esp`: stack pointer**
- Special registers.
 - Control registers `%cr0`, `%cr2`, `%cr3`, `%cr4`;
 - Segment registers `%cs`, `%ds`, `%es`, `%fs`, `%gs`, `%ss`
 - Global and local descriptor table registers `%gdtr`, `%ldtr`
- Other registers not used in xv6: 8 80-bit floating point registers, debug registers

Stack pointers

- Stack grows downwards
- `%ebp` points to return address
- `%esp` points to top of stack

<code>pushl %eax</code>	<code>subl \$4, %esp</code> <code>movl %eax, (%esp)</code>
<code>popl %eax</code>	<code>movl %eax, (%esp)</code> <code>addl \$4, %esp</code>

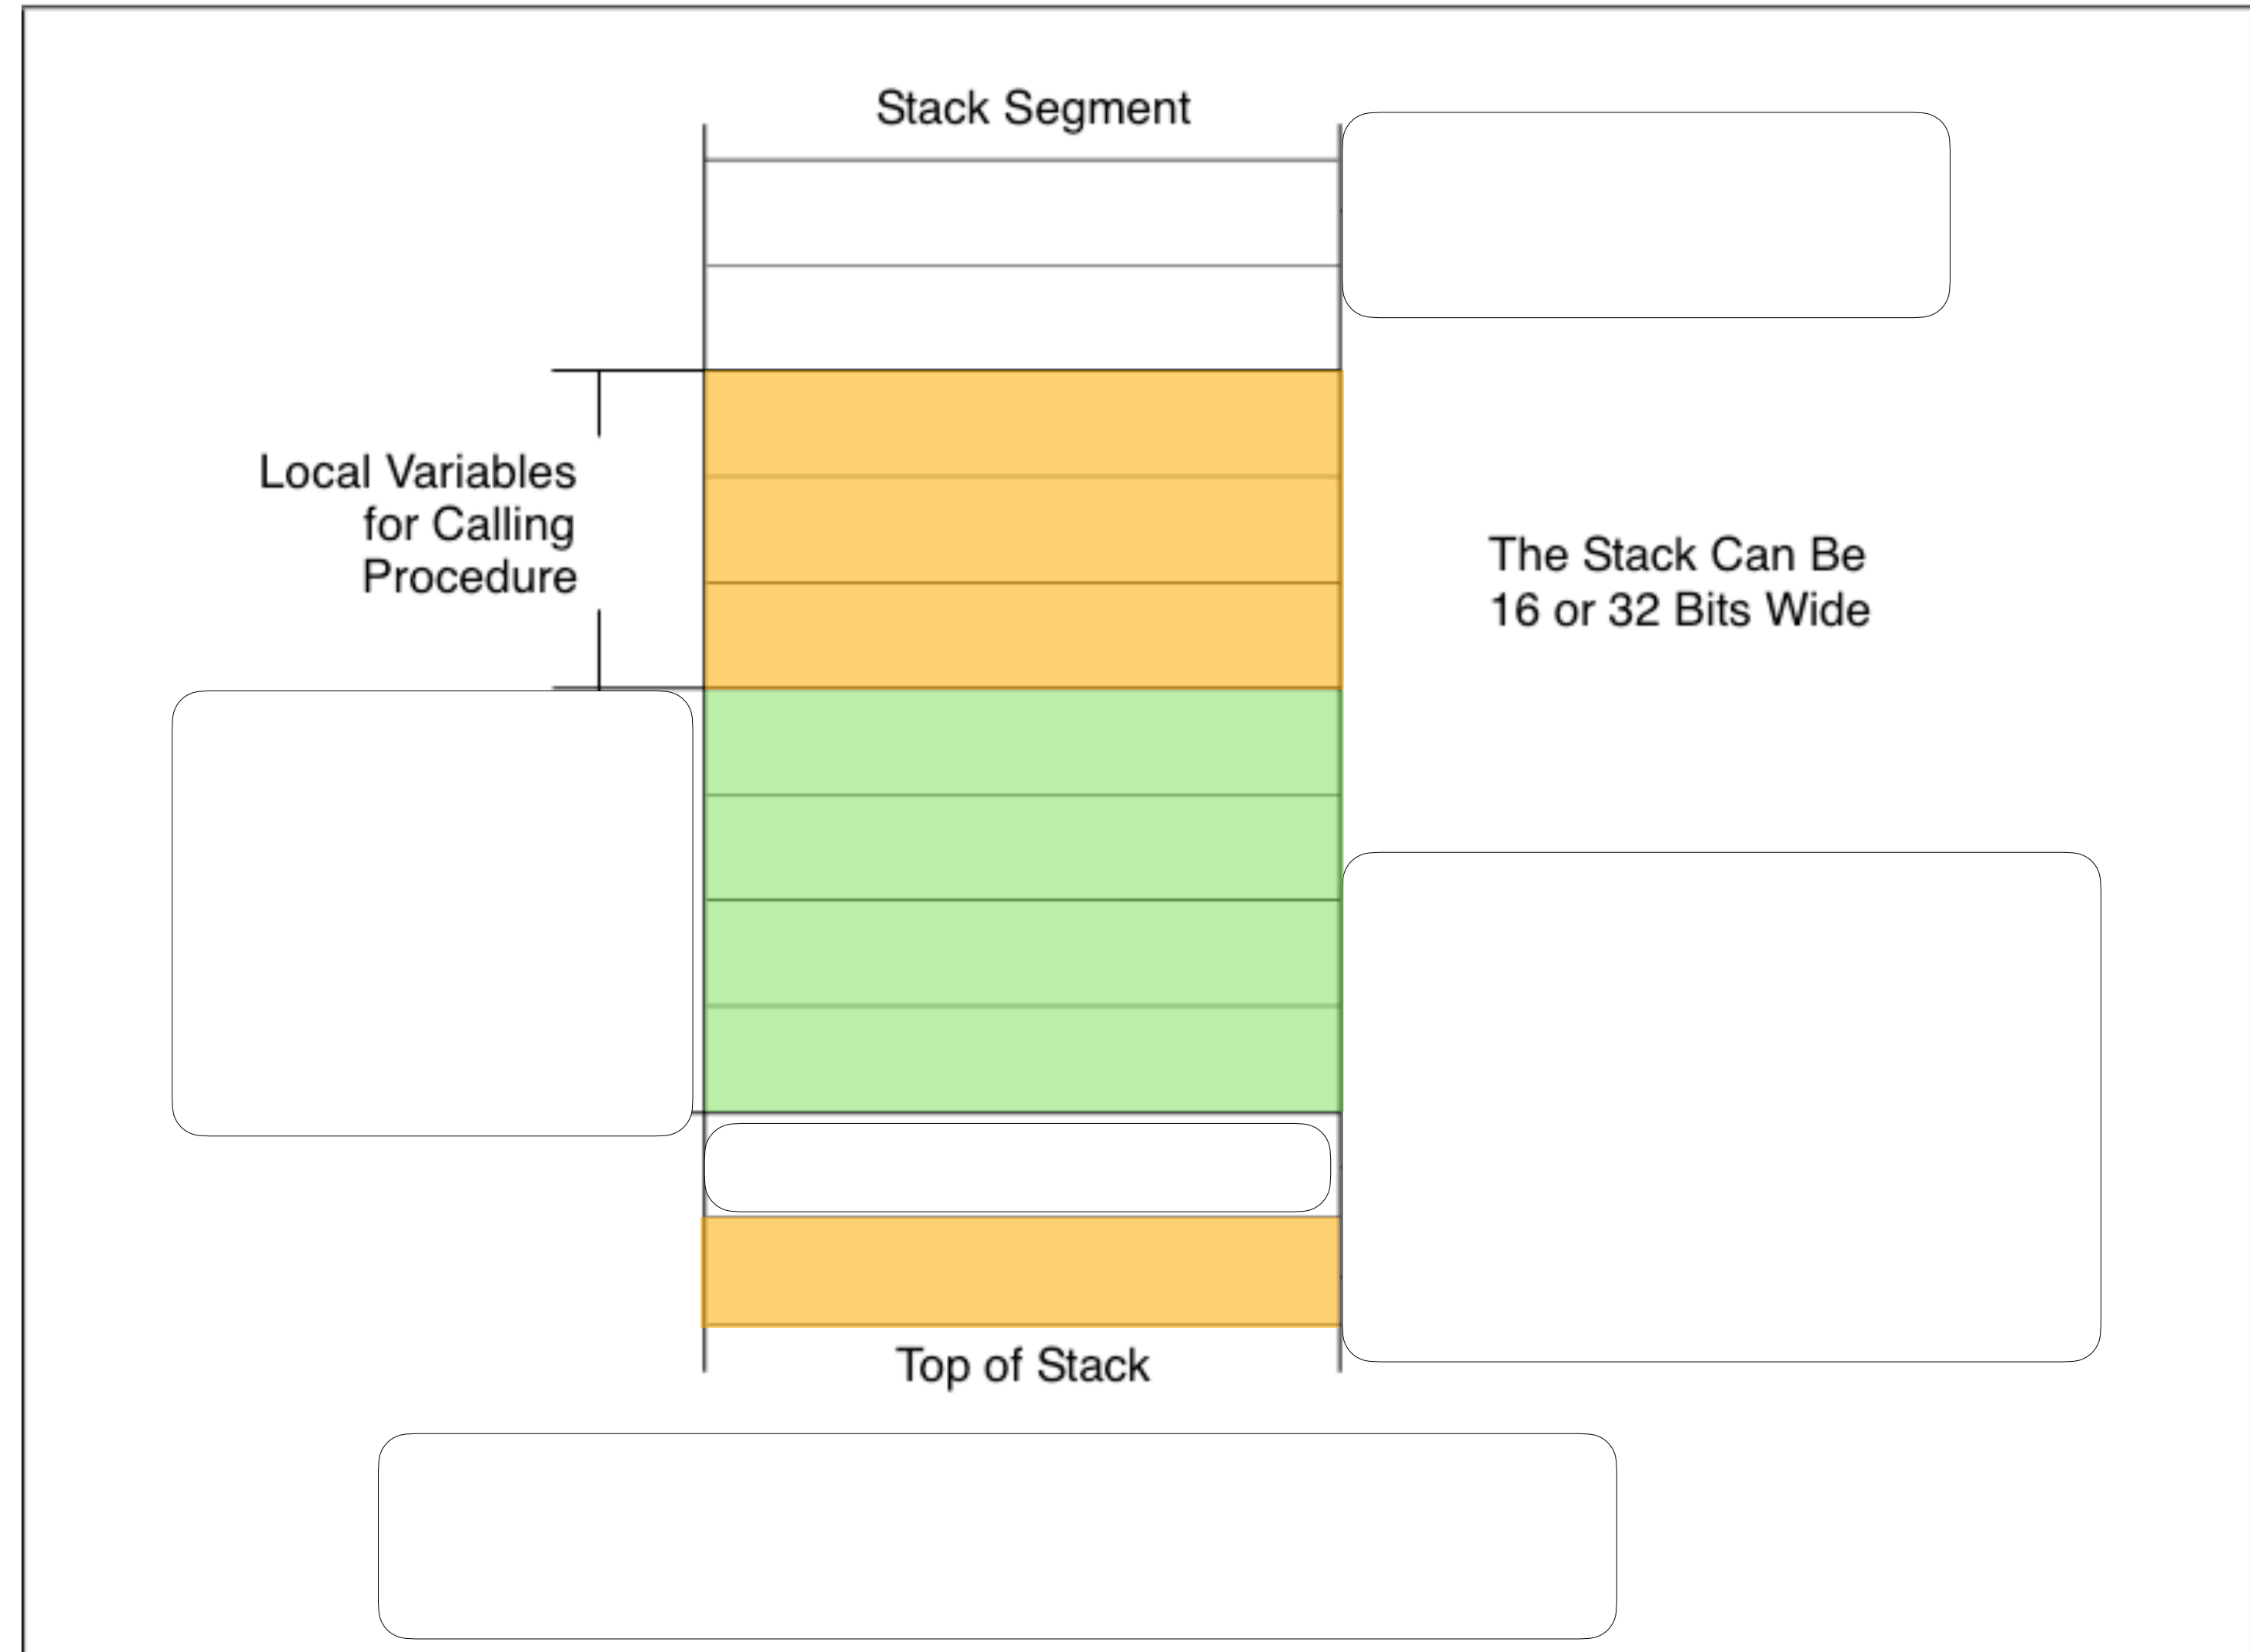


Figure 6-1. Stack Structure

Calling a function

main -> foo(x, y) -> bar(z)

- main pushes foo's parameters (x, y) on the stack
- Executes CALL instruction to save return address on the stack and jump %eip to first instruction of foo
- foo reads parameters from the stack into registers, does computation on them
- foo pushes bar's parameters (z) on the stack, executes CALL instruction
- bar reads z from the stack into registers, does computation on them
- Executes RET instruction to jump %eip to return address in the function foo
- foo executes RET instruction

Function calling in action

02.c

```
int foo(int x, int y) {
    return x + y;
}

int main() {
    return foo(41, 42);
}
```

gcc -m32 -S 02.c

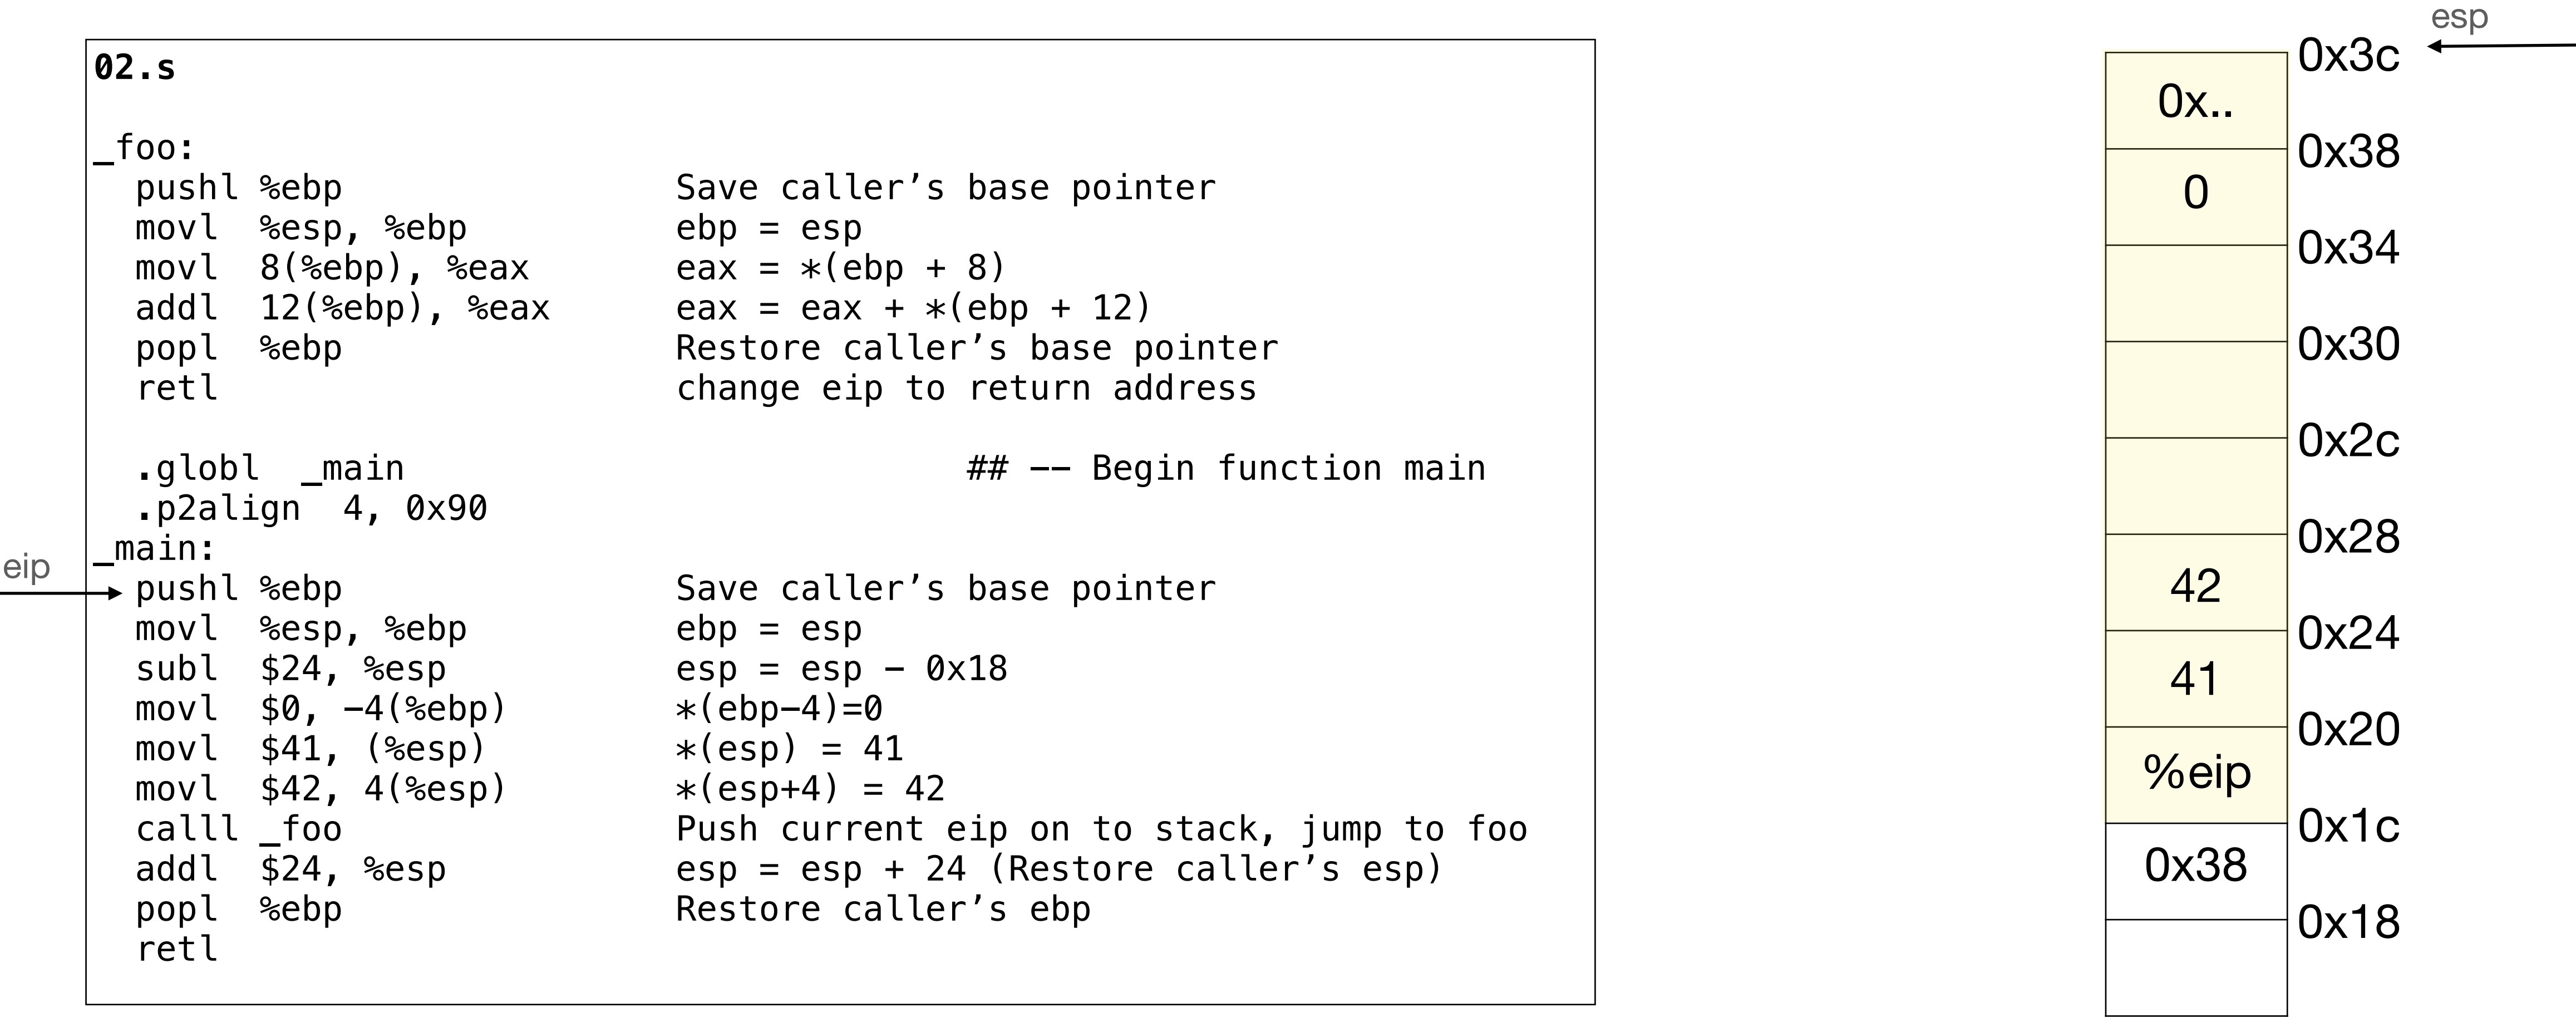
pushl %eax	subl \$4, %esp movl %eax, (%esp)
popl %eax	movl %eax, (%esp) addl \$4, %esp

02.s

```
_foo:
    pushl %ebp                # Save caller's base pointer
    movl  %esp, %ebp         # ebp = esp
    movl  8(%ebp), %eax       # eax = *(ebp + 8)
    addl  12(%ebp), %eax      # eax = eax + *(ebp + 12)
    popl  %ebp               # Restore caller's base pointer
    retl                     # change eip to return address

    .globl _main              ## -- Begin function main
    .p2align 4, 0x90
_main:
    pushl %ebp                # Save caller's base pointer
    movl  %esp, %ebp         # ebp = esp
    subl  $24, %esp          # esp = esp - 24
    movl  $0, -4(%ebp)        # *(ebp-4) = 0
    movl  $41, (%esp)         # *(esp) = 41
    movl  $42, 4(%esp)        # *(esp+4) = 42
    calll _foo                # Push current eip on to stack, jump to foo
    addl  $24, %esp           # esp = esp + 24 (Restore caller's esp)
    popl  %ebp               # Restore caller's ebp
    retl
```

Function calling in action: Stack



gcc calling convention

at entry to a function (i.e. just after call):

- %eip points at first instruction of function
- %esp points at return address
- %esp+4 points at first argument

after ret instruction:

- %eip contains return address
- %esp points at arguments pushed by caller

called function may have trashed arguments

- %eax contains return value (or trash if function is void)
- %eax, %edx, and %ecx may be trashed (caller save)
- %ebp, %ebx, %esi, %edi must contain contents from time of call (callee save)

Instructions are in memory!

02.s

```
_foo:
    pushl %ebp
    movl  %esp, %ebp
    movl  8(%ebp), %eax
    addl  12(%ebp), %eax
    popl  %ebp
    retl

    .globl _main
    .p2align 4, 0x90
_main:
    pushl %ebp
    movl  %esp, %ebp
    subl  $24, %esp
    movl  $0, -4(%ebp)
    movl  $41, (%esp)
    movl  $42, 4(%esp)
    calll _foo
    addl  $24, %esp
    popl  %ebp
    retl
```

```
gcc -m32 -c 02.s -o 02.o
vim 02.o
:%!xxd
```



Instructions are in memory!

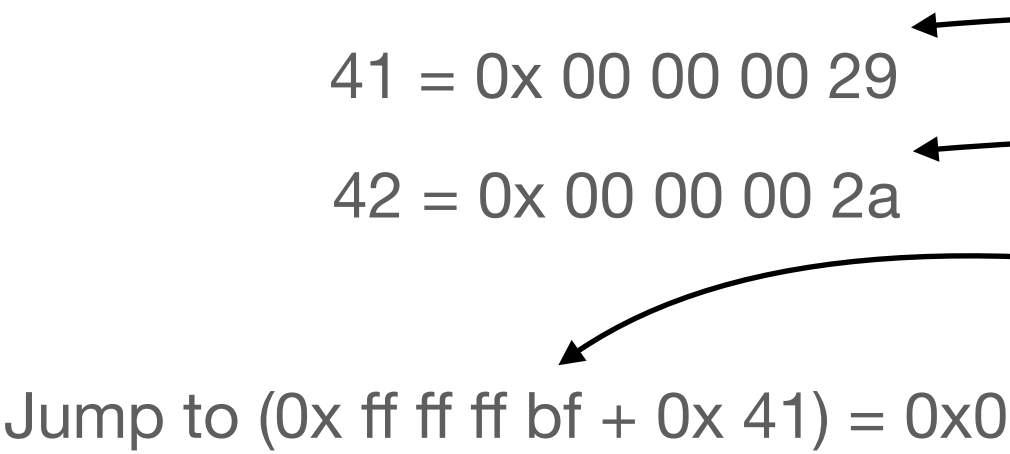
```
02.s

_foo:
    pushl %ebp
    movl  %esp, %ebp
    movl  8(%ebp), %eax
    addl  12(%ebp), %eax
    popl  %ebp
    retl

.globl _main
.p2align 4, 0x90
_main:
    pushl %ebp
    movl  %esp, %ebp
    subl  $24, %esp
    movl  $0, -4(%ebp)
    movl  $41, (%esp)
    movl  $42, 4(%esp)
    calll _foo
    addl  $24, %esp
    popl  %ebp
    retl
```

gcc -m32 -c 02.s -o 02.o
objdump -d 02.o > 02.dump

call 0x0123	pushl %eip (*) movl \$0x123, %eip (*)
ret	popl %eip (*)



00000000 <_foo>:
0: 55 pushl %ebp
1: 89 e5 movl %esp, %ebp
3: 8b 45 0c movl 12(%ebp), %eax
6: 8b 45 08 movl 8(%ebp), %eax
9: 8b 45 08 movl 8(%ebp), %eax
c: 03 45 0c addl 12(%ebp), %eax
f: 5d popl %ebp
10: c3 retl

00000020 <_main>:
20: 55 pushl %ebp
21: 89 e5 movl %esp, %ebp
23: 83 ec 18 subl \$24, %esp
26: c7 45 fc 00 00 00 00 movl \$0, -4(%ebp)
2d: c7 04 24 29 00 00 00 movl \$41, (%esp)
34: c7 44 24 04 2a 00 00 00 movl \$42, 4(%esp)
3c: e8 bf ff ff ff calll 0x0 <_foo>
41: 83 c4 18 addl \$24, %esp
44: 5d popl %ebp
45: c3 retl

* fake instructions
call saves eip of next instruction

Compiling, linking, loading

- *Preprocessor* takes C source code (ASCII text), expands #include etc, produces C source code
- *Compiler* takes C source code (ASCII text), produces assembly language (also ASCII text) *02.main.c -> 02.main.s*
- *Assembler* takes assembly language (ASCII text), produces *.o file* (binary, machine-readable!) *02.main.s -> 02.main.o*
- *Linker* takes multiple *‘.o’s*, produces a single *program image a.out* (binary) *02.main.o, 02.func.o -> 02.main*
- *Loader* loads the program image into memory at run-time and starts executing it

Memory access hierarchy: caches

- Registers are limited in size.
- Main memory is slow.
- Recently accessed data lives on on-chip caches.
- Mostly transparent to OS

Intel Core i7 Xeon 5500 at 2.4 GHz		
Memory	Access time	Size
register	1 cycle	64 bytes
L1 cache	~4 cycles	64 kilobytes
L2 cache	~10 cycles	4 megabytes
L3 cache	~40-75 cycles	8 megabytes
remote L3	~100-300 cycles	
Local DRAM	~60 nsec	
Remote DRAM	~100 nsec	

Figure A-1. Latency numbers for an Intel i7 Xeon system, based on http://software.intel.com/sites/products/collateral/hpc/vtune/performance_analysis_guide.pdf.

I/O devices

Port-mapped IO

- Similar to reading from (writing to) memory locations
- Special instructions:
 - inb (outb) reads (writes) a byte to port
- Only 1024 ports

Writing a byte to line printer

```
#define DATA_PORT    0x378
#define STATUS_PORT   0x379
#define CONTROL_PORT  0x37A
#define    BUSY 0x80
#define    STROBE 0x01
void
lpt_putc(char c)
{
    /* wait for printer to consume previous byte */
    while((inb(STATUS_PORT) & BUSY) == 0);

    /* put the byte on the data lines */
    outb(DATA_PORT, c);

    /* tell the printer to look at the data */
    outb(CONTROL_PORT, STROBE);
    outb(CONTROL_PORT, 0);
}
```

I/O devices

Memory-mapped IO

- Regular memory access instructions
- Reads and writes are routed to appropriate device
 - Writes to VGA memory appear on the screen
- Power-on jumps %eip to 0x000F000
- Careful! Does not behave like memory!
 - Reading same location twice can change due to external events

