

Input output

Polling vs interrupt, Interrupt controllers, interrupt descriptor table, interrupt handlers, Direct memory access, hard disks

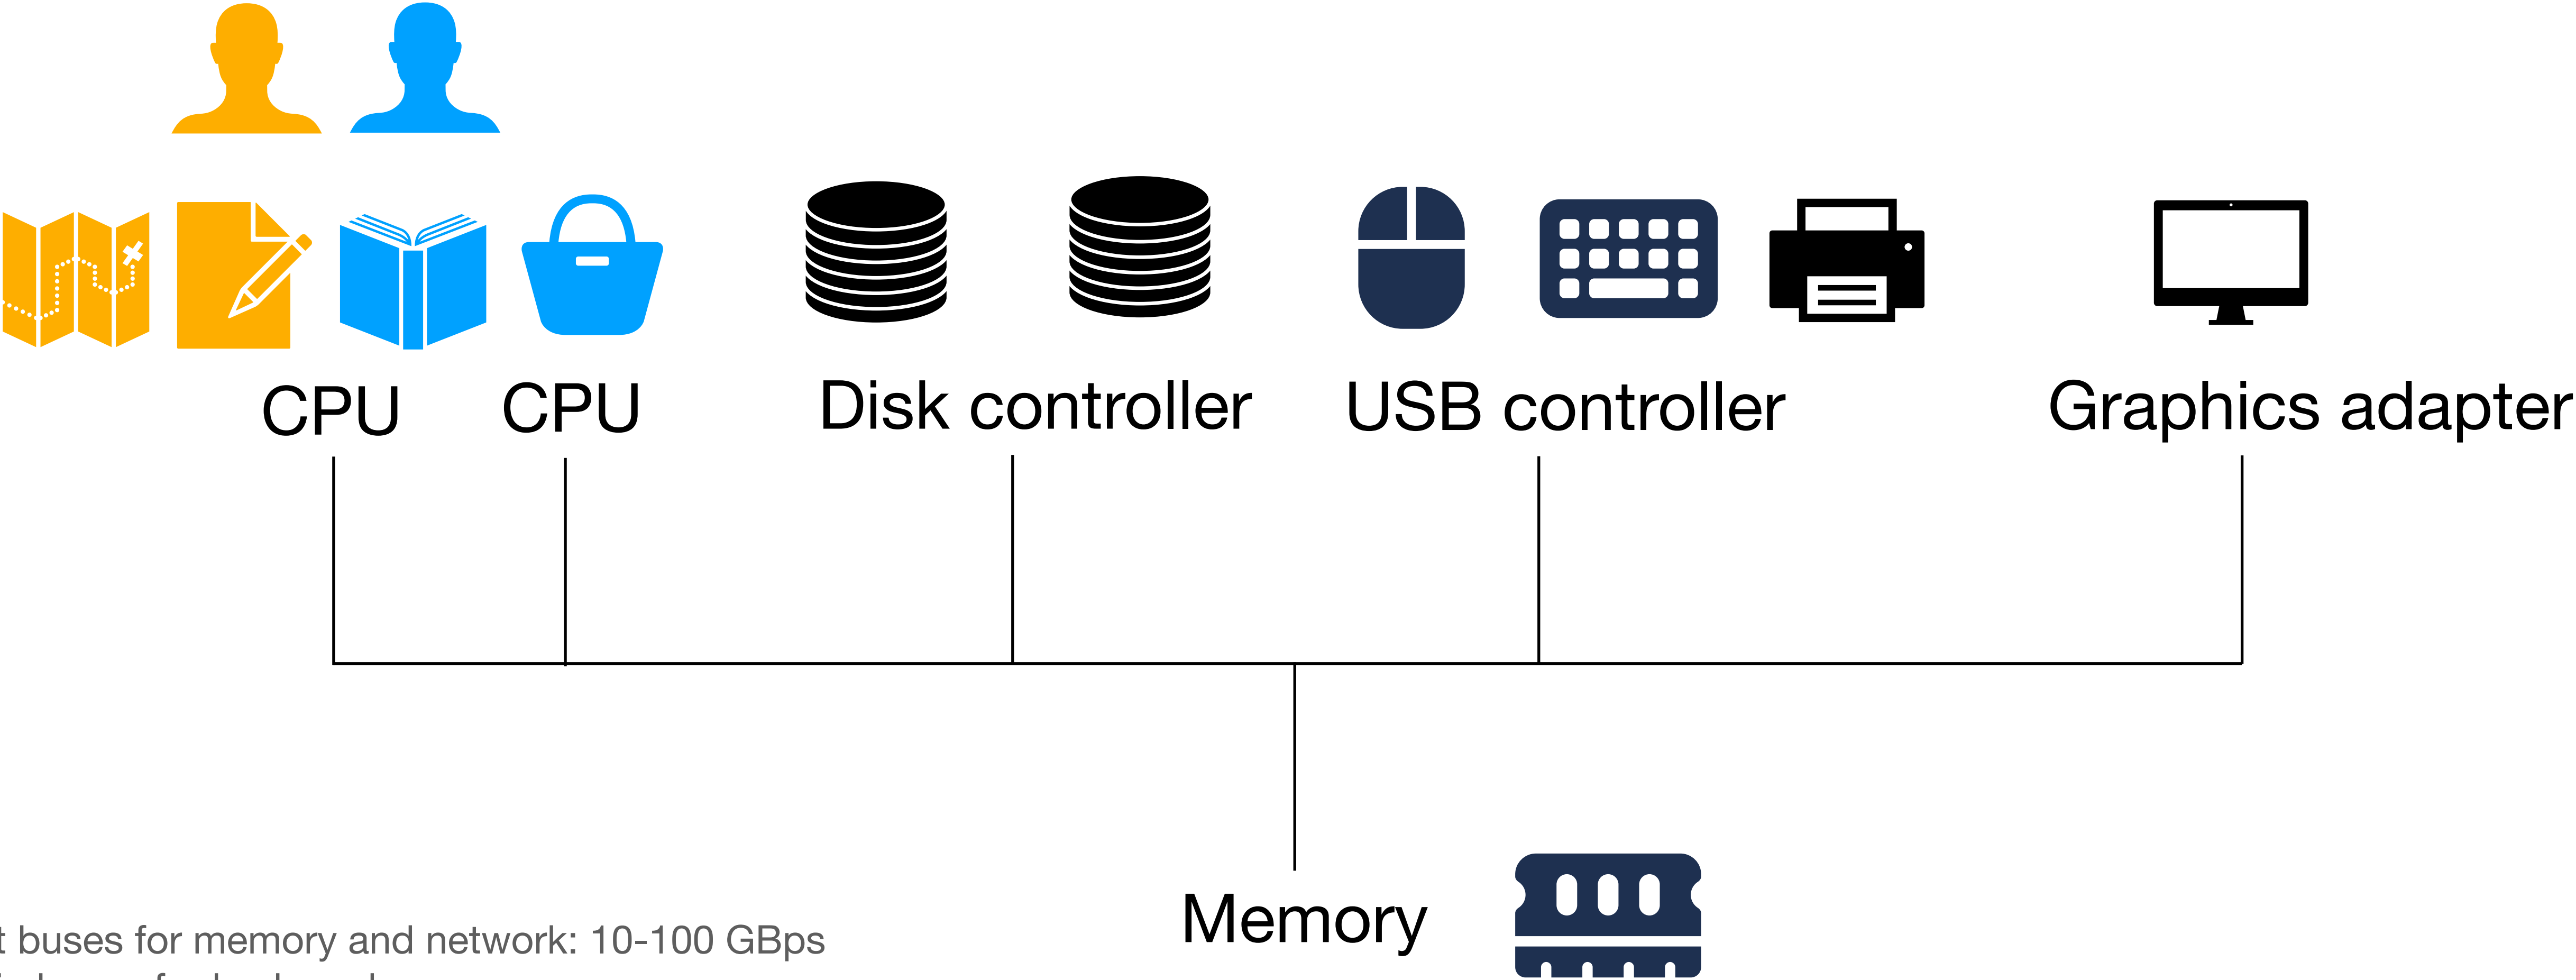
Agenda

- Overview of IO devices (OSTEP Ch. 36): Polling, Interrupts, Direct memory access
- Interrupt handling (xv6 Ch. 3): interrupt controllers, interrupt descriptor table
- Hard disk drives (OSTEP Ch. 37): disk geometry, disk scheduling
- Redundant Array of Inexpensive Disks (OSTEP Ch. 38): improve capacity, throughput, fault tolerance

Overview of IO devices

OSTEP Ch. 36

Computer organization




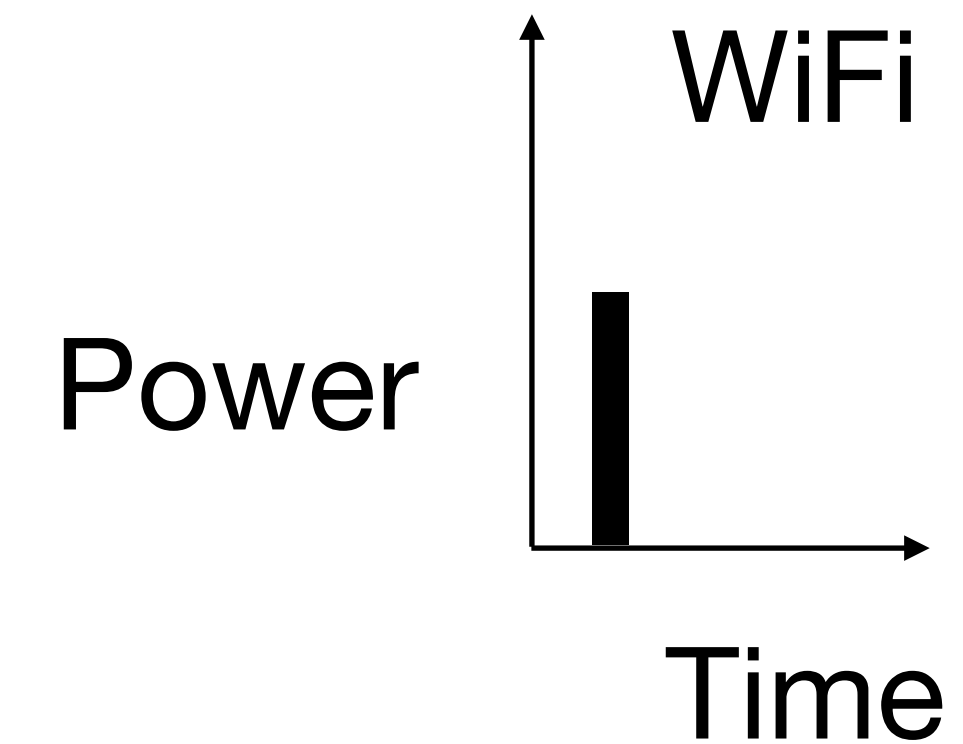


Fat buses for memory and network: 10-100 GBps
Thin buses for keyboard, mouse

Fitting into the OS

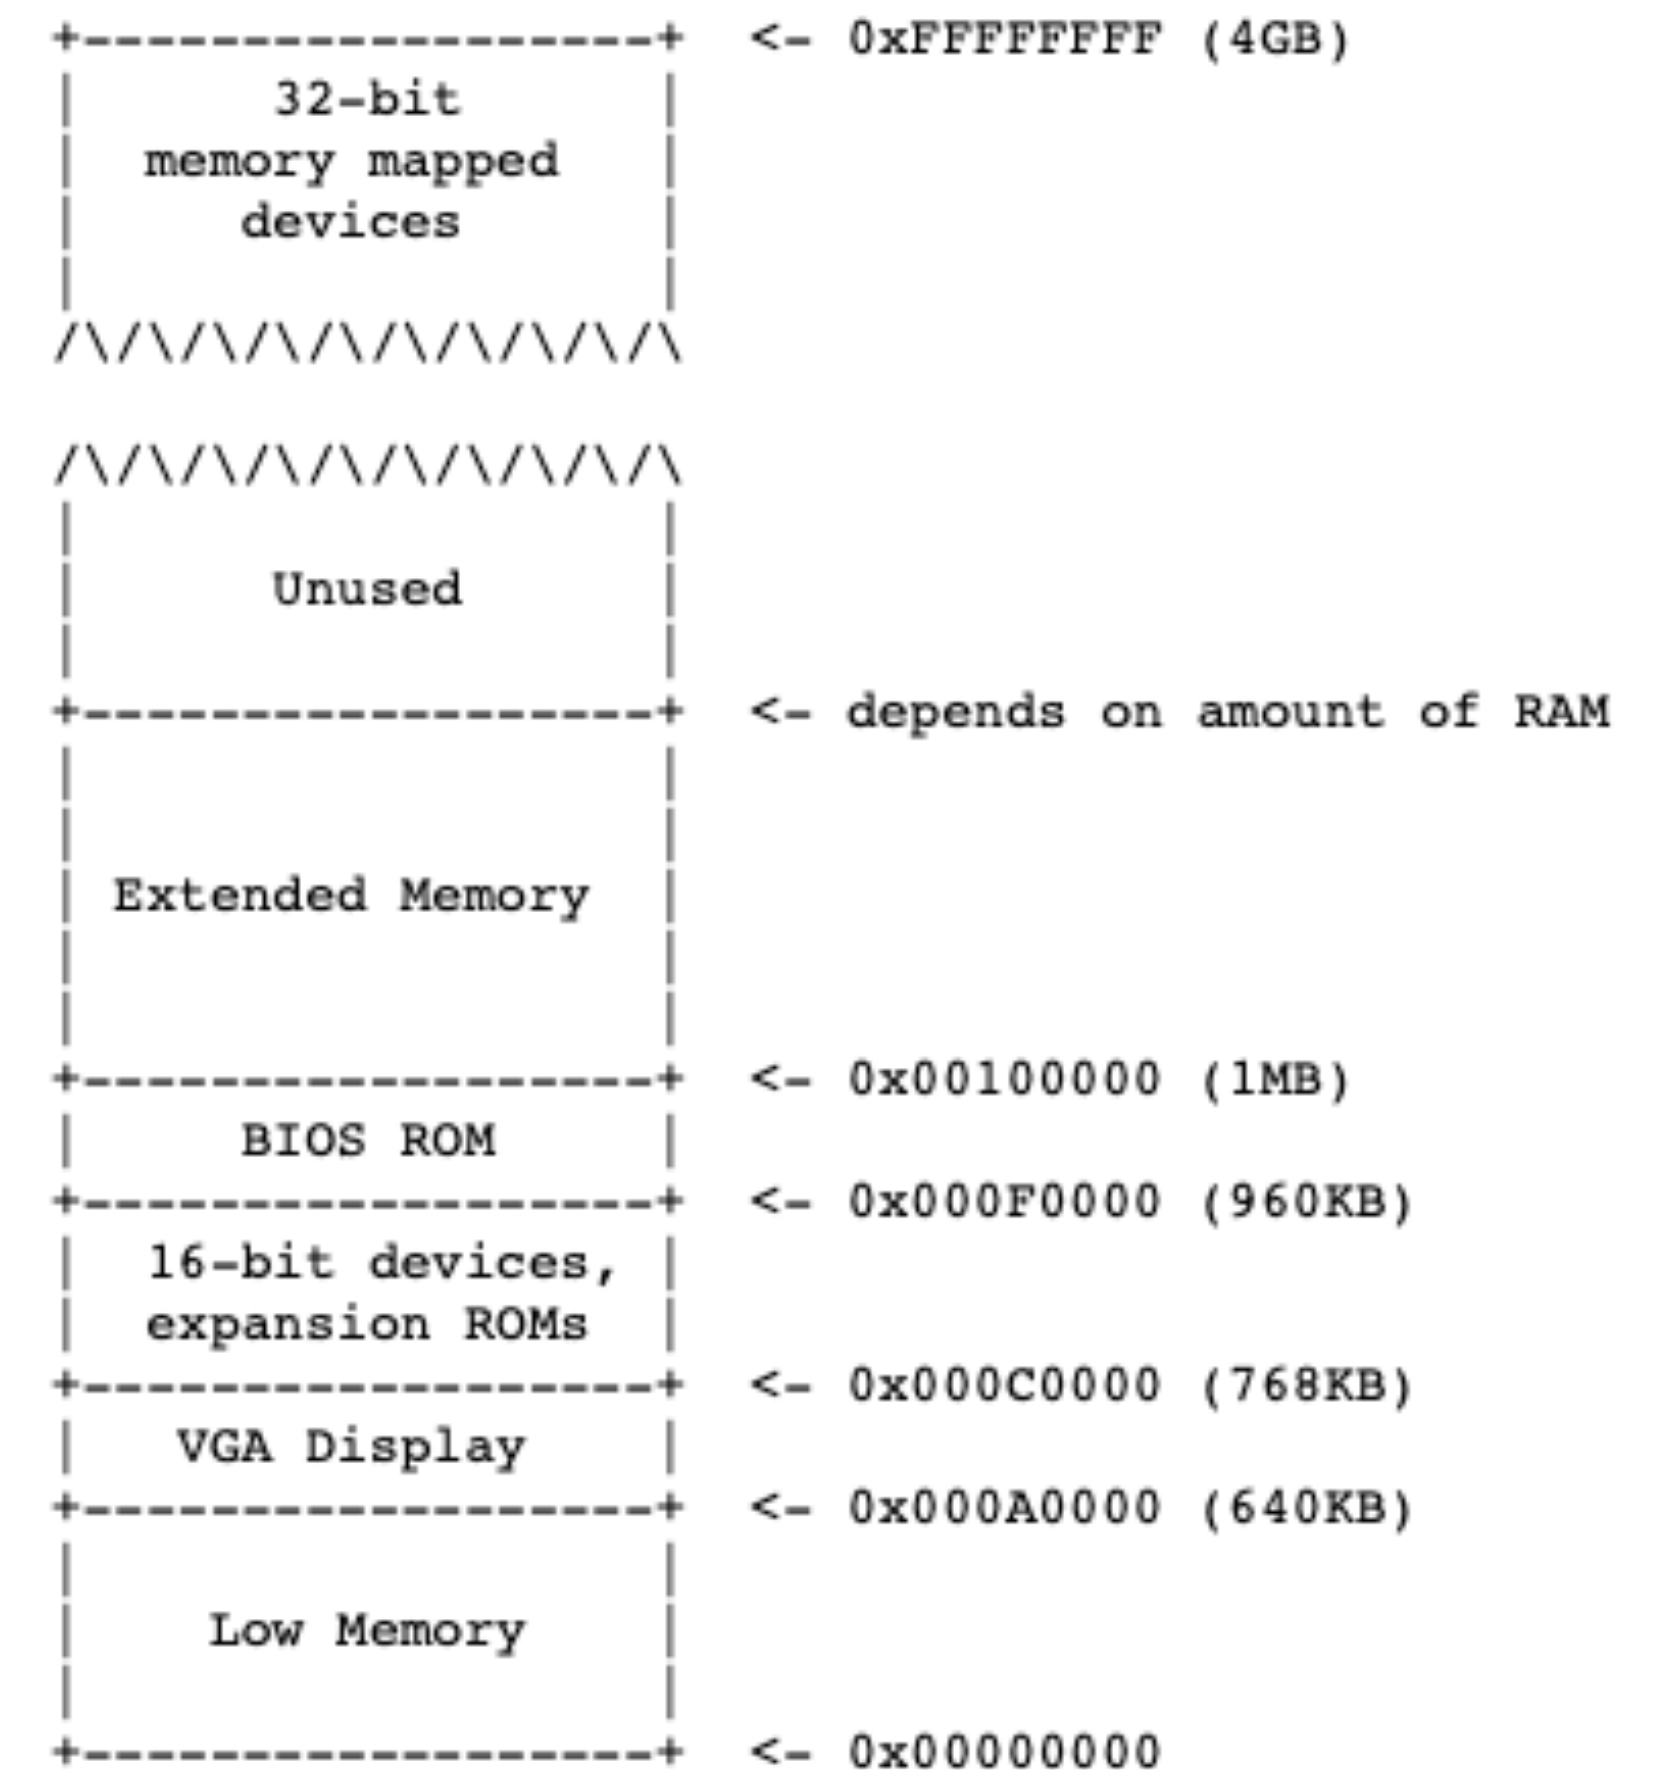
Hide device specific details in device driver

- Abstraction allows OS and applications to stay device-neutral
- Abstraction can hurt.
 - Example: 3G/4G are more efficient for bursty traffic
- > 70% of OS code is device drivers. Tend to have most number of bugs



Memory-mapped IO and Port-mapped IO

- Memory mapped:
 - Regular memory access instructions
 - Reads and writes are routed to appropriate device
 - Does not behave like memory! Reading same location twice can change due to external events
- Port mapped:
 - Special IN and OUT instructions



Canonical protocol

- Poll device until it is ready
- CPU cannot do anything else.
- Example: CPU needs to spend ~1 million instructions waiting for disk
- Ok for bootloader. It does not have anything else to do.
- Not ok for OS. It can run other processes.




Figure 36.3: A Canonical Device

bootmain.c



```
void waitdisk(void){
    // Wait for disk ready.
    while((inb(0x1F7) & 0xC0) != 0x40);
}

// Read a single sector at offset into dst.
void readsect(void *dst, uint offset) {
    // Issue command.
    waitdisk();
    outb(0x1F2, 1);    // count = 1
    ..
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors

    // Read data.
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}
```

Lowering CPU overheads with interrupts

- Device sends an interrupt that it is ready
- CPU runs another process in the meantime
- Better CPU utilisation
- Not a good idea if device is fast.
 - If first poll finds that the device is ready, unnecessary overhead of switching processes



More efficient data movement

Direct Memory Access (DMA)

CPU	1	1	1	2	2	c	c	c	1
Disk				1	1				

CPU	1	1	1	2	2	2	2	2	1
DMA						c	c	c	
Disk				1	1				

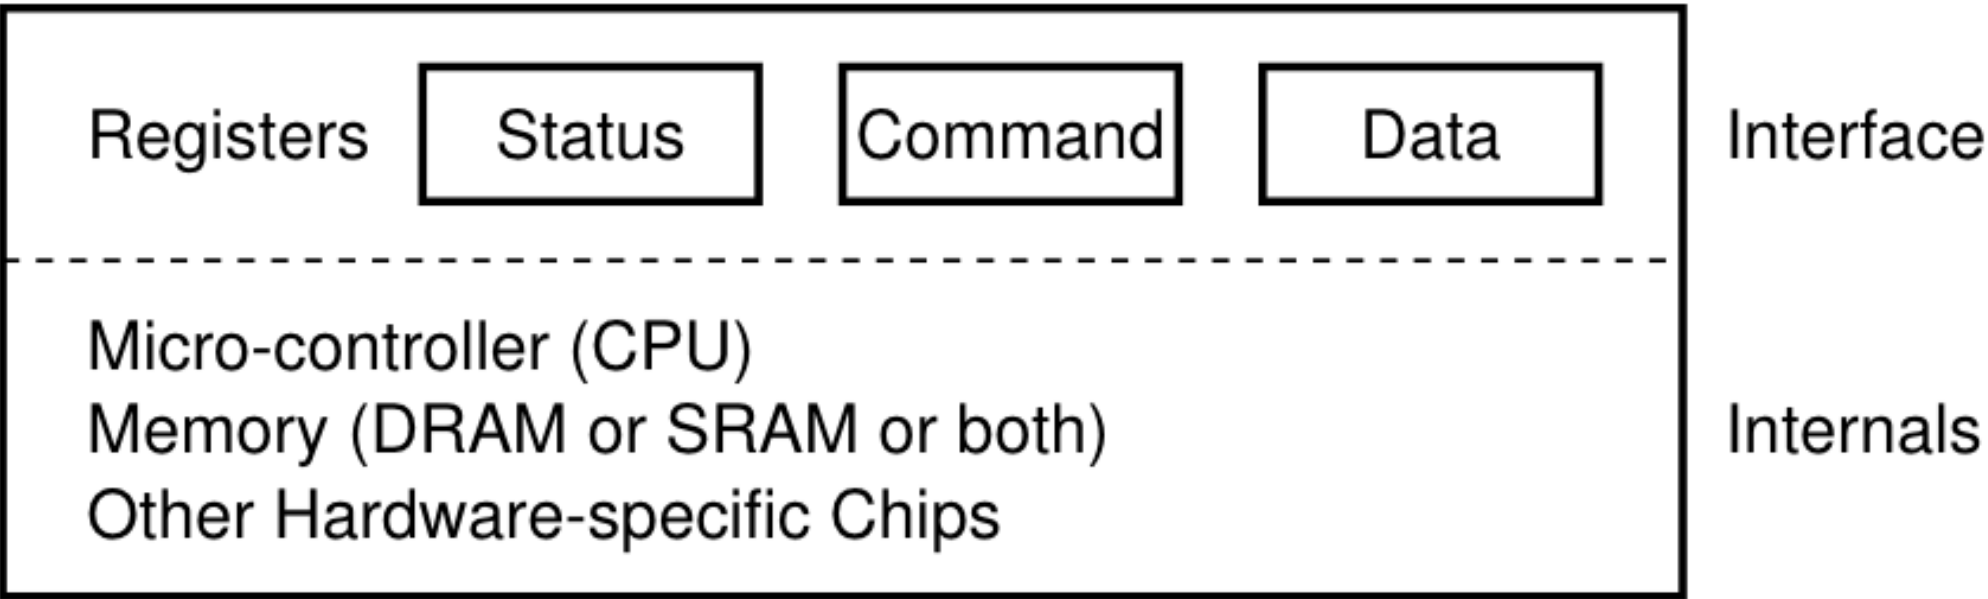


Figure 36.3: A Canonical Device

bootmain.c

```
void waitdisk(void){
    // Wait for disk ready.
    while((inb(0x1F7) & 0xC0) != 0x40);
}

// Read a single sector at offset into dst.
void readsect(void *dst, uint offset) {
    // Issue command.
    waitdisk();
    outb(0x1F2, 1);    // count = 1
    ..
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors

    // Read data.
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}
```

Interrupt controllers, interrupt handling

xv6 Ch. 3 “Code: interrupts”

Calculator analogy



20
10
30
50
30
10
20
10

- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)

Interrupt
←

Give me the calculator!

- $3 * 2 = 6$

End of Interrupt
←

Ok, you can have it back

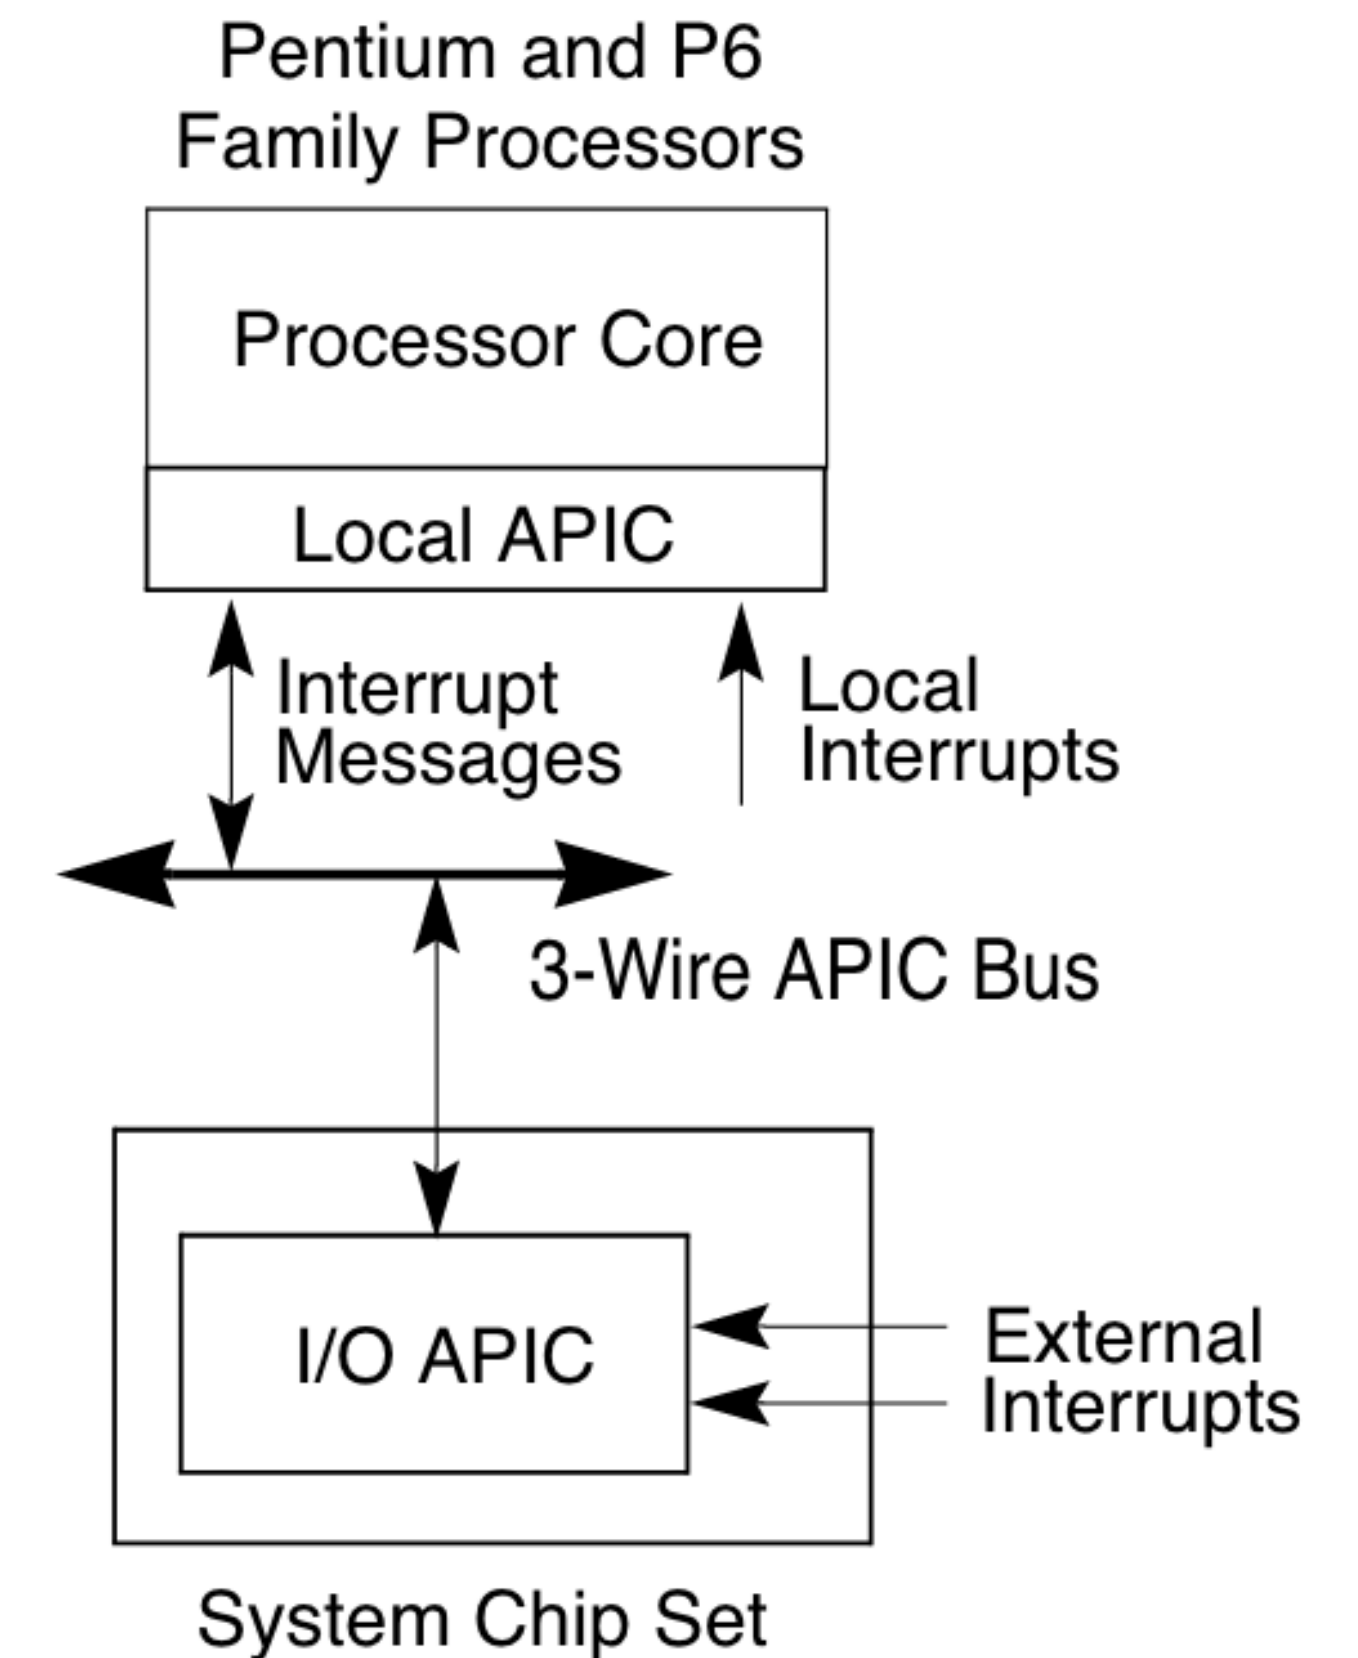
- + 5 0 = (move pointer to 30)
- + 3 0 = (move pointer to 10)
- + 1 0 = (move pointer to 20)
- + 2 0 = (move pointer to 10)



Figure 3b. Interrupt Method

Advanced programmable interrupt controllers (APIC)

- Each CPU can have local APICs for handling *local interrupts* like timer, thermal sensor, etc.
- A separate IO APIC receives external interrupts like keyboard, mouse, disk, etc and forwards it to a particular CPU
- Example: Route keyboard interrupts to CPU-0, disk interrupts to CPU-1



Code walkthrough

- main.c calls lapicinit, picinit, ioapicinit
- lapicinit enables timer interrupt at every 10ms. lapicw is just writing to memory location (MMIO)
- picinit just disables PIC using outb instructions (PMIO)
- ioapicinit initialises IO APIC with MMIO
- Bootloader had disabled interrupt with cli. We will not receive interrupts yet.

Interrupt enable flag

- cli: Clear interrupt flag
 - PIC is not allowed to interrupt
- sti: Set interrupt flag


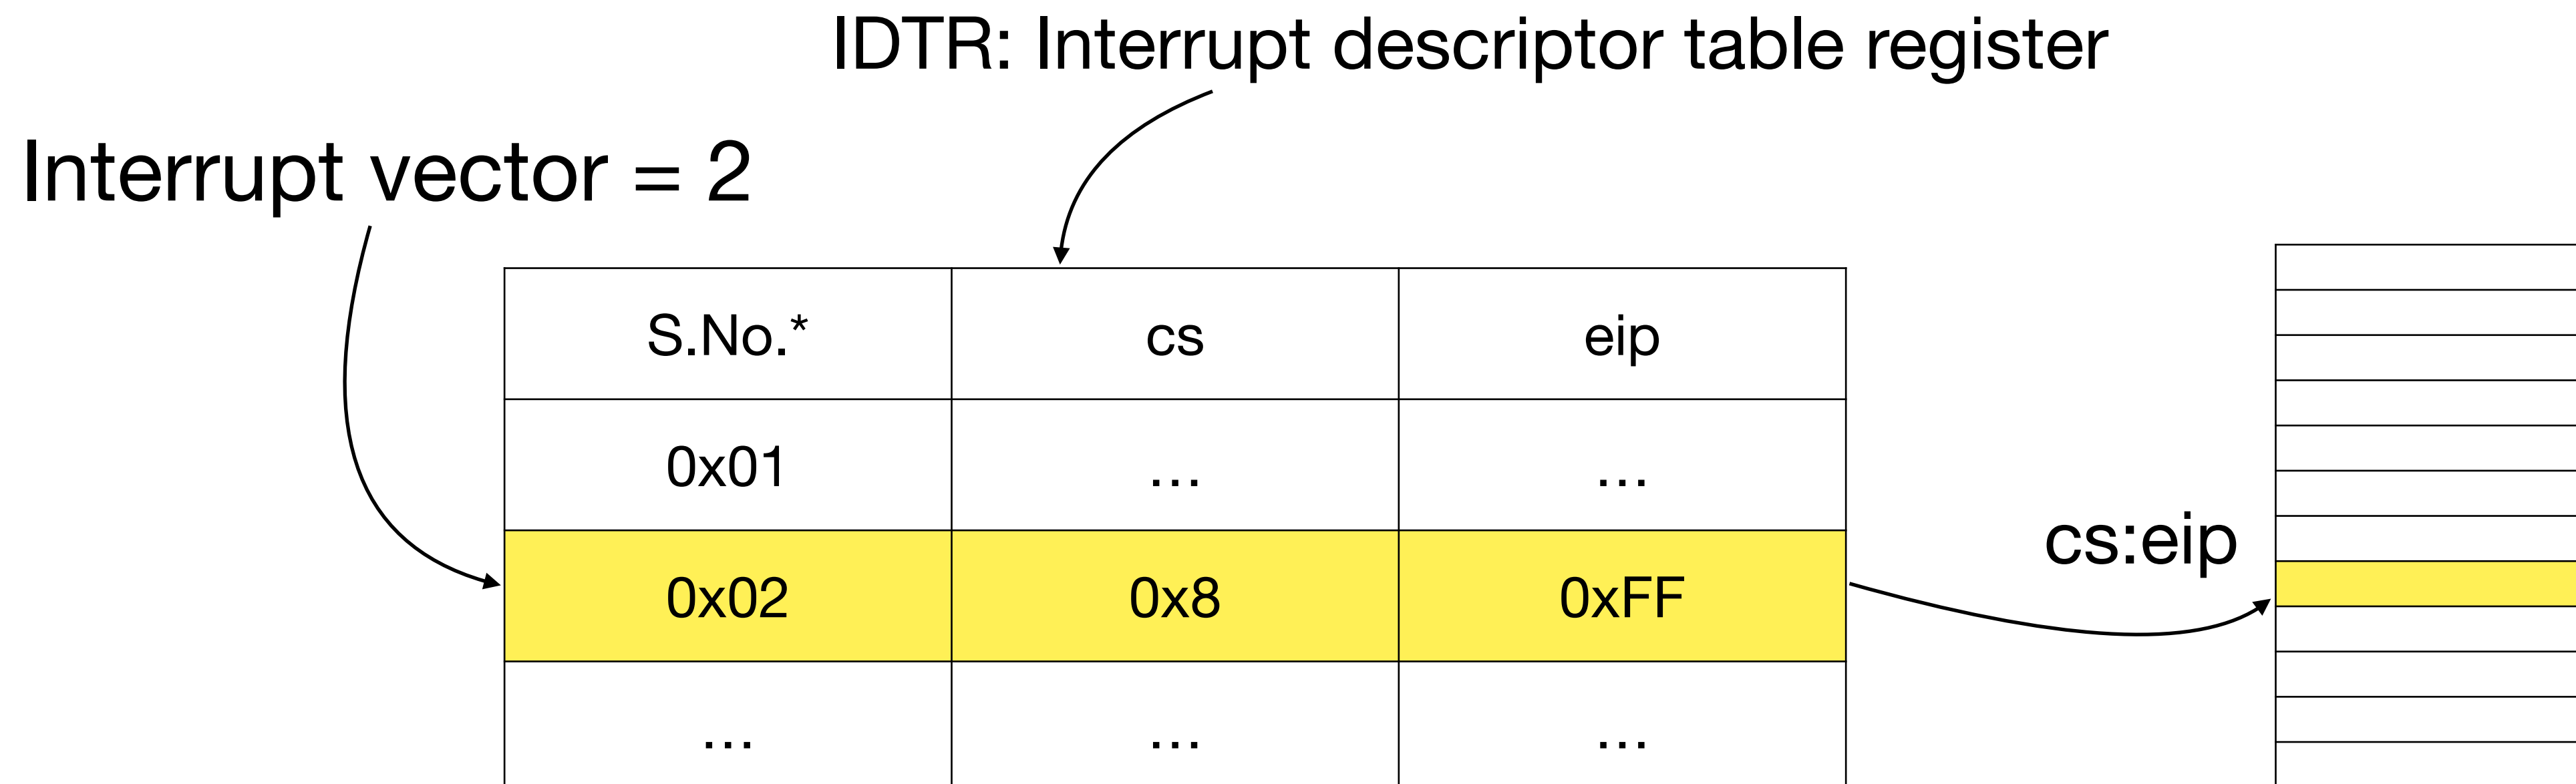


Figure 3-8. EFLAGS Register

Interrupt handling in a nutshell

- OS sets up “interrupt descriptor table” (IDT)
- Points IDTR to IDT using LIDT instruction
- When interrupt occurs, jump %eip to interrupt handler, handle interrupt, tells LAPIC about end of interrupt, resume what we were doing



Interrupt descriptor table

- Interrupt descriptor table register (IDTR) points to interrupt descriptor table in memory
- OS sets up IDT and initialises IDTR using LIDT instruction
- Interrupt descriptor table has one entry for each interrupt vector (upto $2^8=256$)




Figure 6-1. Relationship of the IDTR and IDT

Interrupt descriptor table (2)

- Each IDT entry is 64-bits. Contains **code segment** and **eip**
- When interrupt appears, hardware changes CS and EIP to the one pointed by IDT entry

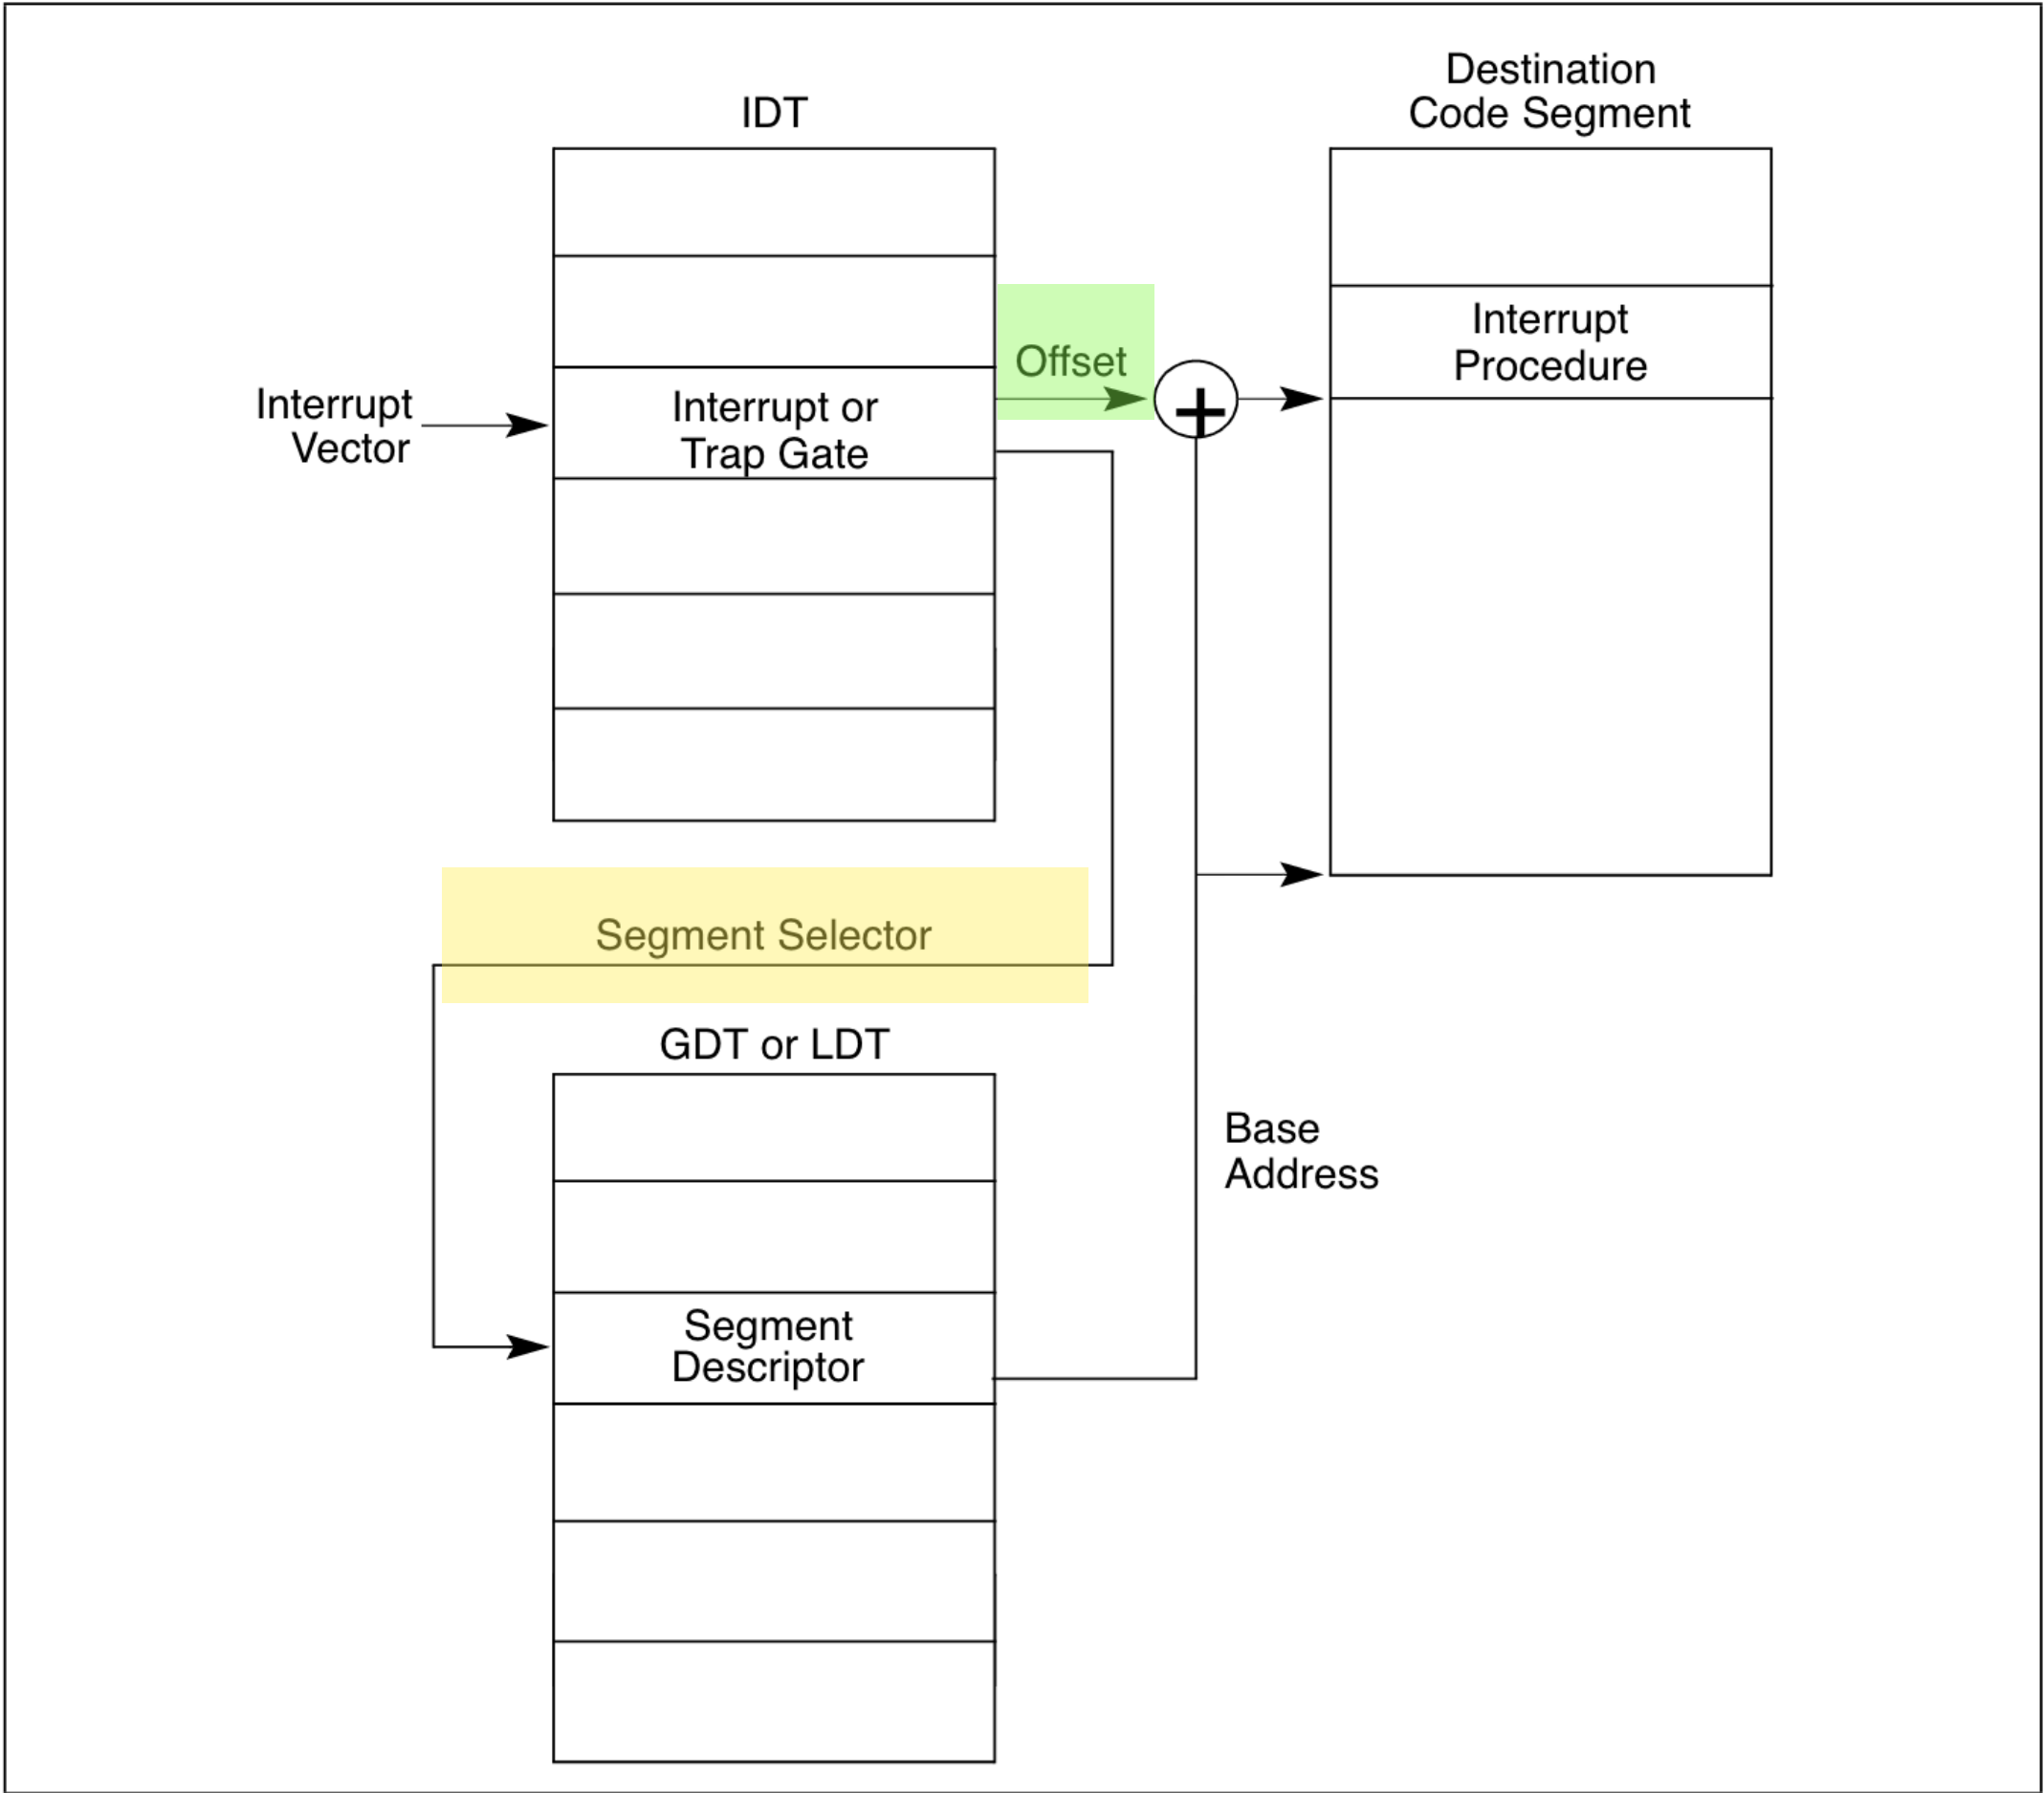
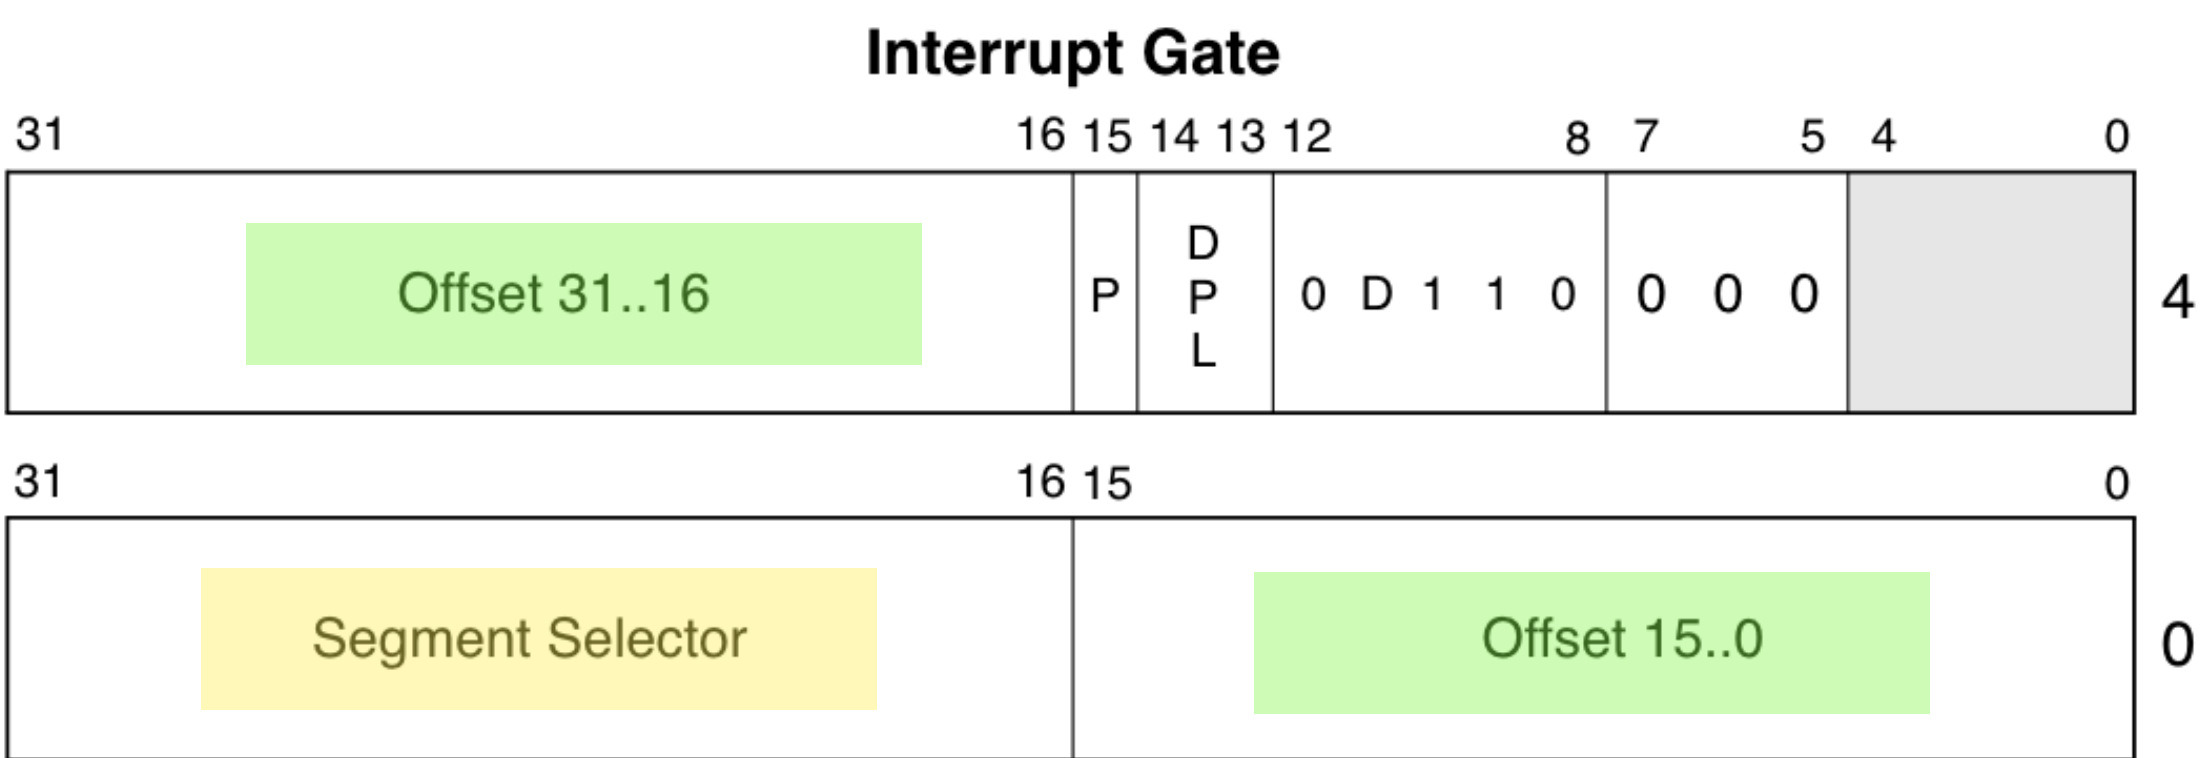
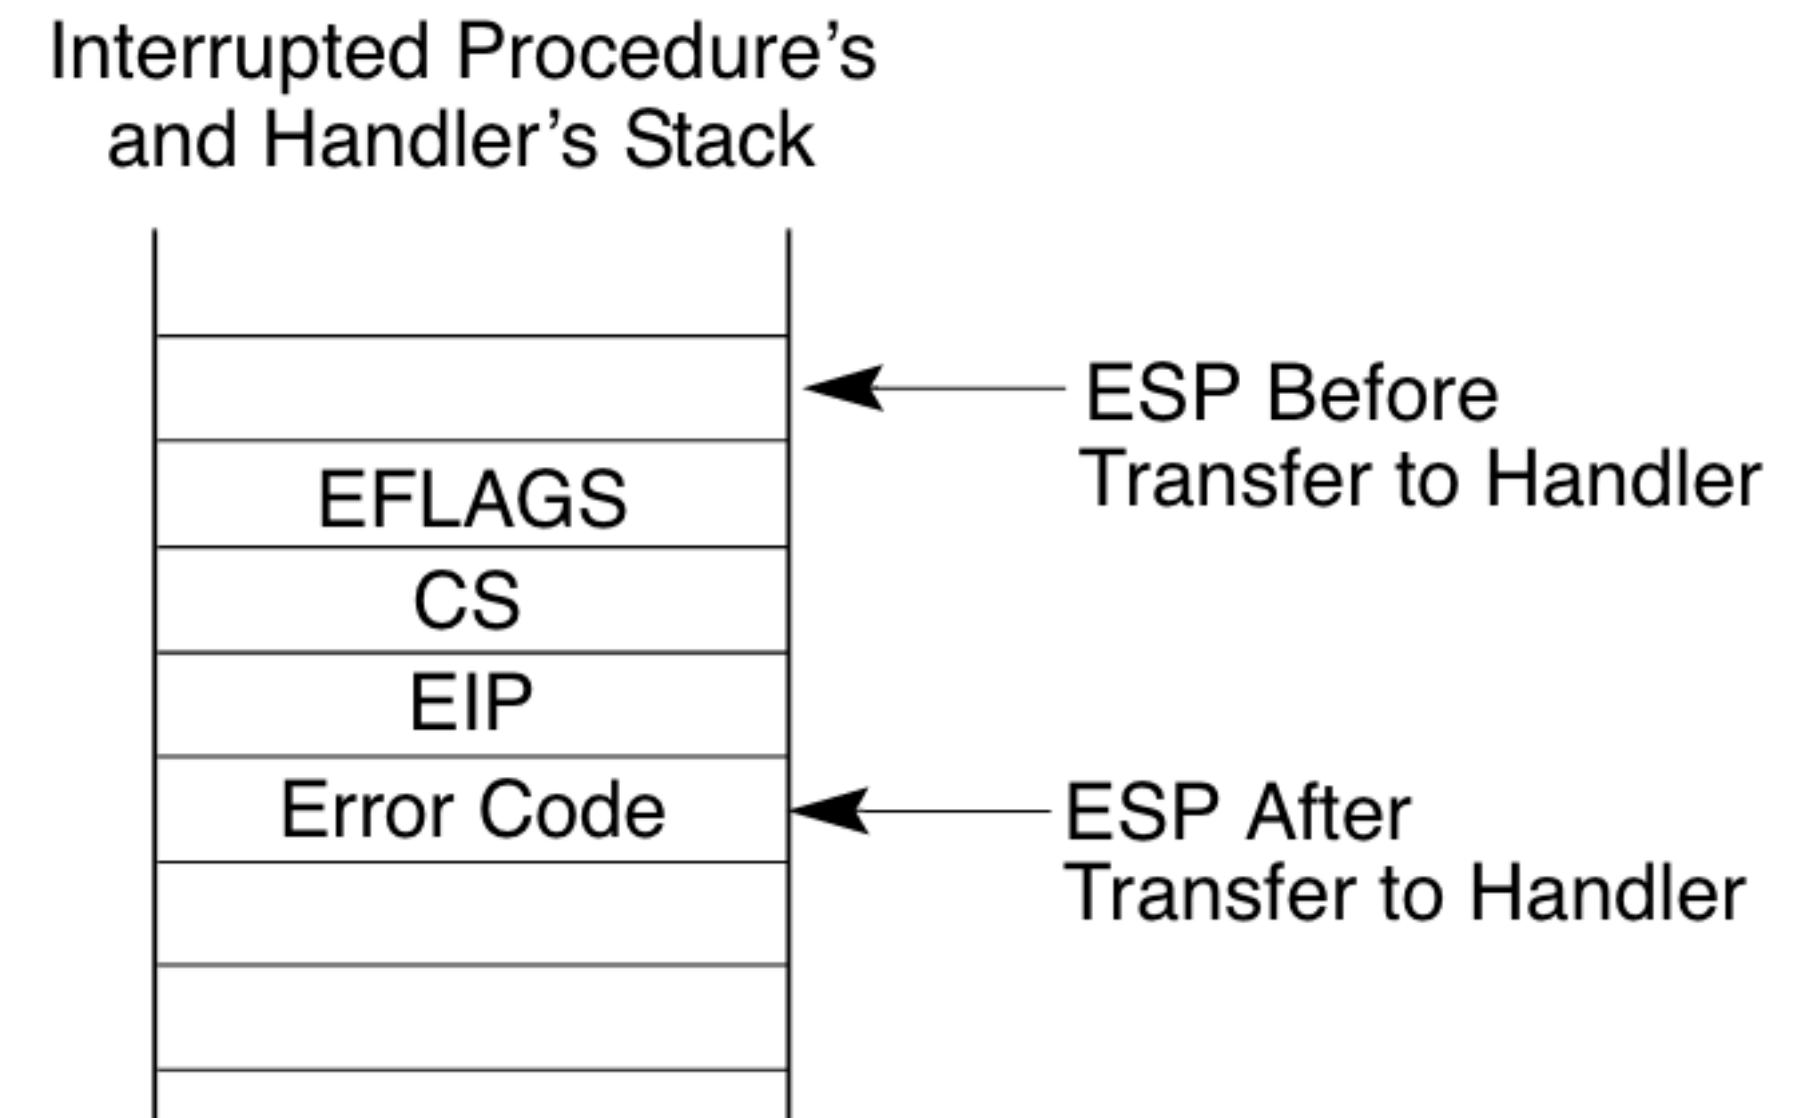


Figure 6-3. Interrupt Procedure Call

Interrupt handling

- On an interrupt, hardware pushes old EFLAGS, CS and EIP on the stack
- Jumps CS and EIP according to IDT
- IRET instruction (similar to RET instruction) restores CS, EIP, EFLAGS, ESP
- Interrupt handler may push more registers, like eax etc. on the stack.



Code walkthrough

p4-traps

- vectors.pl creates 256 IDT entries. 'i'th entry write 'i' on top of the stack and jumps to 'alltraps'
- main.c calls tvinit and idtinit to setup interrupt descriptor table to populate the 256 entries and point IDTR to IDT
- 'alltraps' in trapasm.S runs 'pushal' to save general purpose registers. Then it calls 'trap' with the trapframe.
- 'trap' in 'trapasm.S' reads trapno saved by vectors.S to find out which interrupt occurred. It handles timer and spurious interrupts. It signals EOI to LAPIC when it is done with interrupt.
- trapasm recovers registers with popal, backs up esp above err code and trap number, executes iret to jump back to whatever OS was doing earlier

Visualizing interrupt handling

